

Probabilistic Supervised Learning

Debaditya Roy



Probabilistic Supervised Learning

Goal: Learn the **conditional probability** of output y given input \mathbf{x}

$$p(y|\mathbf{x})$$

Choose distribution based on the type of output:

- Real → Gaussian
- Binary → Bernoulli
- Multiclass → Multinoulli
- Count → Poisson
- Positive real → Gamma

Discriminative vs. Generative Learning

Discriminative Models

- Learn **directly**: $p(y|x)$
- Focus on **prediction**
- Don't model input x

Real output (Gaussian):

$$p(y|x) = N(\mathbf{w}^\top \mathbf{x}, \beta^{-1})$$

Binary output (Bernoulli):

$$p(y|x) = \text{Bernoulli}(\sigma(\mathbf{w}^\top \mathbf{x}))$$

Generative Models

- Learn joint: $p(y|x) = \frac{p(x,y)}{p(x)}$
- Can be used for prediction and data generation
- Require modeling both x and y

Multiclass output (Multinoulli):

$$\begin{aligned} p(\mathbf{x}|y=k) &= \mathcal{N}(\boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k) \\ p(y) &= \text{Categorical} \end{aligned}$$

Count output (Poisson):

$$p(y | x) = \text{Poisson}(\exp(\mathbf{w}^\top \mathbf{x}))$$

Note: Generative approach can also be used for other settings too, such as unsupervised learning and semi-supervised learning

Probabilistic Linear Regression

Goal: We **model uncertainty** – instead of just finding one best line, we define a *probabilistic model* over the data.

Given:

- ▶ Training data $\{(\mathbf{x}_n, y_n)\}_{n=1}^N$, where $\mathbf{x}_n \in \mathbb{R}^D$, $y_n \in \mathbb{R}$
- ▶ Assume:

$$y_n = \mathbf{w}^\top \mathbf{x}_n + \epsilon_n, \quad \epsilon_n \sim \mathcal{N}(0, \beta^{-1})$$

Likelihood:

$$p(y_n \mid \mathbf{x}_n, \mathbf{w}, \beta) = \mathcal{N}(y_n \mid \mathbf{w}^\top \mathbf{x}_n, \beta^{-1})$$

Prior on Weights

Prior Assumption:

$$p(w) = \mathcal{N}(0, \lambda^{-1} I)$$

Interpretation:

- ▶ Zero-mean Gaussian prior over weights
- ▶ Equivalent to L2 regularization
- ▶ λ controls strength of regularization

Remarks:

- ▶ Non-zero mean or structured covariance priors can be used if prior knowledge is available

Posterior over Weights

Bayes' Rule:

$$p(w | X, y) \propto p(y | X, w) \cdot p(w)$$

Closed-form posterior (Gaussian):

$$p(w | X, y) = \mathcal{N}(\mu_N, \Sigma_N)$$

$$\Sigma_N = (\lambda I + \beta X^\top X)^{-1}$$

$$\mu_N = \beta \Sigma_N X^\top y$$

Posterior is Gaussian due to conjugacy.

Derivation of posterior

$$\log p(\mathbf{w} \mid \mathbf{X}, \mathbf{y}) \propto \log p(\mathbf{y} \mid \mathbf{X}, \mathbf{w}) + \log p(\mathbf{w})$$

$$\log p(\mathbf{y} \mid \mathbf{X}, \mathbf{w}) = -\frac{\beta}{2} \underbrace{(\mathbf{y} - \mathbf{X}\mathbf{w})^\top (\mathbf{y} - \mathbf{X}\mathbf{w})}_{(\mathbf{y} - \mathbf{X}\mathbf{w})^\top (\mathbf{y} - \mathbf{X}\mathbf{w}) = \mathbf{y}^\top \mathbf{y} - 2\mathbf{y}^\top \mathbf{X}\mathbf{w} + \mathbf{w}^\top \mathbf{X}^\top \mathbf{X}\mathbf{w}}$$

$$\log p(\mathbf{w}) = -\frac{\lambda}{2} \mathbf{w}^\top \mathbf{w}$$

$$(\mathbf{y} - \mathbf{X}\mathbf{w})^\top (\mathbf{y} - \mathbf{X}\mathbf{w}) = \mathbf{y}^\top \mathbf{y} - 2\mathbf{y}^\top \mathbf{X}\mathbf{w} + \mathbf{w}^\top \mathbf{X}^\top \mathbf{X}\mathbf{w}$$

$$\log p(\mathbf{w} \mid \mathbf{X}, \mathbf{y}) \propto -\frac{1}{2} [\mathbf{w}^\top \underbrace{(\beta \mathbf{X}^\top \mathbf{X} + \lambda \mathbf{I}) \mathbf{w}}_{\mathbf{w}^\top A \mathbf{w}} - 2\beta \underbrace{\mathbf{y}^\top \mathbf{X}\mathbf{w}}_{\mathbf{b}^\top \mathbf{w}}]$$

$$\mathbf{w}^\top A \mathbf{w} - 2\mathbf{b}^\top \mathbf{w} = (\mathbf{w} - \mathbf{w}_0)^\top A(\mathbf{w} - \mathbf{w}_0) - \mathbf{w}_0^\top A \mathbf{w}_0$$

$$\mathbf{w}_0 = A^{-1} \mathbf{b} = (\beta \mathbf{X}^\top \mathbf{X} + \lambda \mathbf{I})^{-1} (\beta \mathbf{X}^\top \mathbf{y})$$

$$\mu_N = A^{-1} \mathbf{b} = \beta \Sigma_N \mathbf{X}^\top \mathbf{y}$$

$$\Sigma_N = (\lambda \mathbf{I} + \beta \mathbf{X}^\top \mathbf{X})^{-1}$$

$$p(y_n \mid \mathbf{x}_n, \mathbf{w}, \beta) = \mathcal{N}(y_n \mid \mathbf{w}^\top \mathbf{x}_n, \beta^{-1})$$


$$p(\mathbf{w}) = \mathcal{N}(0, \lambda^{-1} \mathbf{I})$$

Posterior: A Visualization


- Assume a lin. reg. problem with true $\mathbf{w} = [w_0, w_1]$, $w_0 = -0.3$, $w_1 = 0.5$
- Assume data generated by a linear regression model $y = w_0 + w_1 x + \text{"noise"}$
 - Note: It's actually 1-D regression (w_0 is just a bias term), or 2-D reg. with feature $[1, x]$
- Figures below show the “data space” and posterior of \mathbf{w} for different number of observations (note: with no observations, the posterior = prior)

Each red line represents the “data” generated for a randomly drawn \mathbf{w} from the current posterior

data space



Posterior



Posterior Predictive Distribution

For a new input \mathbf{x}^* :

$$p(y^* \mid \mathbf{x}^*, \mathcal{D}) = \int p(y^* \mid \mathbf{x}^*, \mathbf{w}) p(\mathbf{w} \mid \mathcal{D}) d\mathbf{w}$$

Result: Gaussian distribution


$$p(y^* \mid \mathbf{x}^*, \mathcal{D}) = \mathcal{N}(\boldsymbol{\mu}_N^\top \mathbf{x}^*, \ \boldsymbol{\beta}^{-1} + \mathbf{x}^{*\top} \boldsymbol{\Sigma}_N \mathbf{x}^*)$$

Notes:

- ▶ Mean: $\boldsymbol{\mu}_N^\top \mathbf{x}^*$
- ▶ Variance includes noise + uncertainty in \mathbf{w}
- ▶ More informative than MLE/MAP point estimates


Posterior Predictive Distribution: An Illustration

- Black dots are training examples



- Width of the shaded region at any x denotes the predictive uncertainty at that x (+/- one std-dev)
- Regions with more training examples have smaller predictive variance

Nonlinear Regression



- Can extend the linear regression model to handle nonlinear regression problems
- One way is to replace the feature vectors \mathbf{x} by a nonlinear mapping $\phi(\mathbf{x})$

$$p(y|\mathbf{x}, \mathbf{w}) = \mathcal{N}(\mathbf{w}^\top \phi(\mathbf{x}), \beta^{-1})$$

Can be pre-defined (e.g., replace a scalar x by polynomial mapping $[1, x, x^2]$) or extracted by a pretrained deep neural net

- Alternatively, a **kernel function** can be used to implicitly define the nonlinear mapping
- More on nonlinear regression when we discuss **Gaussian Processes**


Probabilistic Binary Logistic Regression

For $y \in \{0, 1\}$, model:

$$p(y = 1 \mid \mathbf{x}, \mathbf{w}) = \sigma(\mathbf{w}^\top \mathbf{x}) = \frac{1}{1 + e^{-\mathbf{w}^\top \mathbf{x}}}$$

$$p(y \mid \mathbf{x}, \mathbf{w}) = \sigma(\mathbf{w}^\top \mathbf{x})^y \cdot (1 - \sigma(\mathbf{w}^\top \mathbf{x}))^{1-y}$$

Decision rule: Predict 1 if probability > 0.5



Likelihood and MAP

Likelihood over all data:

$$p(\mathbf{y} \mid \mathbf{X}, \mathbf{w}) = \prod_{n=1}^N \sigma(\mathbf{w}^\top \mathbf{x}_n)^{y_n} (1 - \sigma(\mathbf{w}^\top \mathbf{x}_n))^{1-y_n}$$

Negative log-likelihood (Binary Cross-Entropy):

$$\mathcal{L}(\mathbf{w}) = - \sum_{n=1}^N [y_n \log \mu_n + (1 - y_n) \log(1 - \mu_n)]$$

With prior:

$$p(\mathbf{w}) = \mathcal{N}(\mathbf{0}, \lambda^{-1} \mathbf{I}) \propto \exp\left(-\frac{\lambda}{2} \|\mathbf{w}\|^2\right)$$

Encourages smaller weights (regularization)

$$\mathbf{w}_{\text{MAP}} = \arg \min_{\mathbf{w}} \mathcal{L}(\mathbf{w}) + \frac{\lambda}{2} \|\mathbf{w}\|^2$$

Posterior Distribution

Posterior is intractable:

$$p(\mathbf{w} \mid \mathcal{D}) \propto p(\mathbf{w}) \cdot \prod_n \sigma(\mathbf{w}^\top \mathbf{x}_n)^{y_n} (1 - \sigma(\mathbf{w}^\top \mathbf{x}_n))^{1-y_n}$$

- Sigmoid likelihood is not conjugate to the Gaussian prior.
- Posterior does not have a closed-form solution — you can't simplify the product analytically.

Use Laplace Approximation:


$$p(\mathbf{w} \mid \mathcal{D}) \approx \mathcal{N}(\mathbf{w}_{\text{MAP}}, \Lambda^{-1})$$

Laplace's (or Gaussian) Approximation

- Consider a posterior distribution that is intractable to compute

$$p(\theta|\mathcal{D}) = \frac{p(\mathcal{D}, \theta)}{p(\mathcal{D})} = \frac{p(\mathcal{D}|\theta)p(\theta)}{p(\mathcal{D})}$$

- Laplace approximation approximates the above using a **Gaussian** distribution



$$p(\theta|\mathcal{D}) \approx \mathcal{N}(\theta|\theta_{MAP}, \Lambda^{-1})$$

$$\theta_{MAP} = \operatorname{argmax}_{\theta} \log p(\theta|\mathcal{D})$$

$$\Lambda = -\nabla_{\theta}^2 \log p(\theta|\mathcal{D}) \Big|_{\theta=\theta_{MAP}} = -\nabla_{\theta}^2 \log p(\mathcal{D}, \theta) \Big|_{\theta=\theta_{MAP}}$$

Tells us about the space
(curvature) of the true
posterior around θ_{MAP}

Negative of the Hessian,
i.e., the second derivative
of the log joint, at θ_{MAP}

- Laplace's approx. is based on a second-order Taylor approx. of the posterior

Derivation of the Laplace's Approximation

- Let's write the Bayes rule as

$$p(\mathcal{D}) \approx \exp(\log p(\mathcal{D}, \theta_{MAP})) \times (2\pi)^{D/2} \det(\Lambda)^{1/2}$$

$$p(\theta|\mathcal{D}) = \frac{p(\mathcal{D}, \theta)}{p(\mathcal{D})} = \frac{p(\mathcal{D}, \theta)}{\int p(\mathcal{D}, \theta) d\theta} = \frac{\exp[\log p(\mathcal{D}, \theta)]}{\int \exp[\log p(\mathcal{D}, \theta)] d\theta}$$

We also get a Laplace approximation of the marginal likelihood (for free!)

Note: Sometimes marginal likelihood is also called model evidence

- Consider second-order Taylor approximation of a function $f(\theta)$ around some θ_0

$$f(\theta) \approx f(\theta_0) + (\theta - \theta_0)^\top \nabla_\theta f(\theta_0) + \frac{1}{2} (\theta - \theta_0)^\top \nabla_\theta^2 f(\theta_0) (\theta - \theta_0)$$


- Assuming $f(\theta) = \log p(\mathcal{D}, \theta)$ and $\theta_0 = \theta_{MAP}$

$$\log p(\mathcal{D}, \theta) \approx \log p(\mathcal{D}, \theta_{MAP}) + \frac{1}{2} (\theta - \theta_{MAP})^\top \nabla_\theta^2 \log p(\mathcal{D}, \theta_{MAP}) (\theta - \theta_{MAP})$$


$$\begin{aligned} p(\theta|\mathcal{D}) &\propto \exp \left[-\frac{1}{2} (\theta - \theta_{MAP})^\top (-\nabla_\theta^2 \log p(\mathcal{D}, \theta_{MAP})) (\theta - \theta_{MAP}) \right] \\ &= \mathcal{N}(\theta | \theta_{MAP}, \Lambda^{-1}) \quad (\text{where } \Lambda = -\nabla_\theta^2 \log p(\mathcal{D}, \theta_{MAP}) = -\mathbf{H}) \end{aligned}$$

Properties of Laplace's Approximation

- Straightforward if posterior's derivatives (first/second) can be computed easily
- Expensive if parameter θ is very high dimensional
 - Reason: We need to compute and invert Hessian of size $D \times D$ (D is the # of params)
- Can do badly if the (true) posterior is multimodal



For multimodal posteriors,
can use a mixture of
Laplace approximations*



Useful for deep
learning models

If K local modes, then define the approx.
posterior as a mixture of K Gaussians
$$p(\theta|D) \approx \sum_{k=1}^K \pi^{(k)} \mathcal{N}(\theta|\theta_{MAP}^{(k)}, H^{(k)-1})$$

(see paper cited below for details)

- Used only when θ is a real-valued vector (because of Gaussian approximation)
- Note: Even if we have a non-probabilistic model (loss function + regularization), we can obtain an approx. “posterior” for that model using the Laplace's approximation
 - Optima of the regularized loss function will be Gaussian's mean
 - Inverse of the second derivative of the regularized loss function will be covariance matrix

*Mixtures of Laplace Approximations for Improved Post-Hoc Uncertainty in Deep Learning (Eschenhagen et al, 2021)

Posterior Predictive Distribution

- The posterior predictive distribution can be computed as

$$p(y_* = 1 | \mathbf{x}_*, \mathbf{X}, \mathbf{y}) = \int p(y_* = 1 | \mathbf{w}, \mathbf{x}_*) p(\mathbf{w} | \mathbf{X}, \mathbf{y}) d\mathbf{w}$$

Integral not tractable and must be approximated

sigmoid

Gaussian (if using Laplace approx.)

- Monte-Carlo approximation of this integral is one possible way

- Draw M samples $\mathbf{w}_1, \mathbf{w}_2, \dots, \mathbf{w}_M$, from the approx. of posterior
- Approximate the PPD as follows

$$p(y_* = 1 | \mathbf{x}_*, \mathbf{X}, \mathbf{y}) \approx \frac{1}{M} \sum_{m=1}^M p(y_* = 1 | \mathbf{w}_m, \mathbf{x}_*) = \frac{1}{M} \sum_{m=1}^M \sigma(\mathbf{w}_m^\top \mathbf{x}_n)$$

- In contrast, when using MLE/MAP solution $\hat{\mathbf{w}}_{opt}$, the plug-in pred. distribution

$$\begin{aligned} p(y_* = 1 | \mathbf{x}_*, \mathbf{X}, \mathbf{y}) &= \int p(y_* = 1 | \mathbf{w}, \mathbf{x}_*) p(\mathbf{w} | \mathbf{X}, \mathbf{y}) d\mathbf{w} \\ &\approx p(y_* = 1 | \hat{\mathbf{w}}_{opt}, \mathbf{x}_*) = \sigma(\hat{\mathbf{w}}_{opt}^\top \mathbf{x}_n) \end{aligned}$$

Multiclass Logistic Regression

Now, let's extend this to multiple classes $y \in \{1, 2, \dots, K\}$. We model the probability of each class using the softmax function:

$$P(y = k \mid \mathbf{x}, \mathbf{W}) = \frac{\exp(\mathbf{w}_k^\top \mathbf{x})}{\sum_{j=1}^K \exp(\mathbf{w}_j^\top \mathbf{x})}$$

Where:

- $\mathbf{W} = [\mathbf{w}_1, \dots, \mathbf{w}_K]$ are the weight vectors for each class.
- The model outputs a **categorical distribution** over the classes.

MLR: Prior, Likelihood, and Posterior

1. Prior over weights:

$$p(\mathbf{W}) = \prod_{k=1}^K \mathcal{N}(\mathbf{w}_k \mid \mathbf{0}, \tau^{-1} \mathbf{I})$$

2. Likelihood:

Given data $\mathcal{D} = \{(\mathbf{x}_n, y_n)\}$, the likelihood is:

$$p(\mathcal{D} \mid \mathbf{W}) = \prod_{n=1}^N \prod_{k=1}^K P(y_n = k \mid \mathbf{x}_n, \mathbf{W})^{\mathbb{I}(y_n=k)}$$

3. Posterior:

$$p(\mathbf{W} \mid \mathcal{D}) \propto p(\mathbf{W}) \cdot p(\mathcal{D} \mid \mathbf{W})$$

As with binary logistic regression, the posterior is **intractable** due to the softmax in the likelihood. So we approximate it.

MLR: Posterior Distribution

We can approximate $p(\mathbf{W} \mid \mathcal{D})$ using:

- **Laplace Approximation:**
 - Approximate the posterior as Gaussian around the MAP estimate:
$$p(\mathbf{W} \mid \mathcal{D}) \approx \mathcal{N}(\mathbf{W}_{\text{MAP}}, \Lambda^{-1})$$
 - Λ is the **Hessian** of the negative log posterior at the MAP.
- **MCMC Methods:**
 - Sample from the posterior using methods like **Hamiltonian Monte Carlo, Gibbs sampling, or Metropolis-Hastings**.

MLR: Posterior Predictive Distribution

To make predictions for a new point \mathbf{x}_* , we marginalize over the posterior:

$$p(y_* \mid \mathbf{x}_*, \mathcal{D}) = \int p(y_* \mid \mathbf{x}_*, \mathbf{W}) p(\mathbf{W} \mid \mathcal{D}) d\mathbf{W}$$

Since this integral is intractable, we use:

- Monte Carlo approximation:
 - Draw samples $\mathbf{W}^{(s)} \sim p(\mathbf{W} \mid \mathcal{D})$
 - Compute:

$$p(y_* \mid \mathbf{x}_*, \mathcal{D}) \approx \frac{1}{S} \sum_{s=1}^S p(y_* \mid \mathbf{x}_*, \mathbf{W}^{(s)})$$

Generative Supervised Learning

- The conditional distribution $p(y|x)$ can also be defined as

$$p(y|x) = \frac{p(x,y)}{p(x)}$$

- Generative sup. learning is usually more work because $p(x, y)$ has to be estimated
- However, there are some benefits as well. For example, for classification

$p(y)$ is called the “class-prior” or
“class-marginal” distribution

Can incorporate knowledge of frequency
("size") of each class in training data

Can incorporate knowledge of the distribution
("shape") of each class in training data

Can assume simple/sophisticated types
of distributions for the “class-conditional”
distribution $p(x|y)$ and learned them
using the training data of each class

$$p(y|x) = \frac{p(x,y)}{p(x)} = \frac{p(y)p(x|y)}{p(x)}$$

Estimating Class Marginals

- Estimating class marginals $p(y = k)$ is usually straightforward
- Since labels are discrete, we assume class marginal $p(y)$ to be a multinoulli

If only two classes,
assume Bernoulli

$$p(y|\boldsymbol{\pi}) = \text{multinoulli}(y|\pi_1, \pi_2, \dots, \pi_K) = \prod_{k=1}^K \pi_k^{\mathbb{I}[y=k]}$$

$\pi_k = p(y = k)$

These probabilities sum to 1: $\sum_{k=1}^K \pi_k = 1$

- Given N i.i.d. labelled examples $\{(x_n, y_n)\}_{n=1}^N$, $y_n \in \{1, 2, \dots, K\}$ the MLE soln.

$$\boldsymbol{\pi}_{MLE} = \operatorname{argmax}_{\boldsymbol{\pi}} \sum_{n=1}^N \log p(y_n|\boldsymbol{\pi})$$

Subject to constraint $\sum_{k=1}^K \pi_k = 1$

- MLE solution is $p(y = k) = \pi_k = N_k/N$ where $N_k = \sum_{n=1}^N \mathbb{I}[y = k]$
 - Thus $p(y = k) = \pi_k$ is simply the fraction of inputs from class k
- Can also compute MAP estimate or full posterior of $\boldsymbol{\pi}$ using a Dirichlet prior

Estimating Class-Conditionals

To be estimated using the
 N_k training inputs
 $\{\mathbf{x}_n : y_n = k\}$ from class k

- Can assume a distribution $p(\mathbf{x}|y = k) = p(\mathbf{x}|\theta_k)$ for inputs of each class k
- If \mathbf{x} is D -dimensional, $p(\mathbf{x}|\theta_k)$ will be a D -dimensional distribution
- Can compute MLE/MAP estimate or full posterior of θ_k
 - This essentially is a [density estimation](#) problem for the class-cond.
 - In principle, can use any density estimation method
- Choice of the form of $p(\mathbf{x}|\theta_k)$ depends on various factors
 - Nature of input features, e.g.,
 - If $\mathbf{x} \in \mathbb{R}^D$, can use a D -dim Gaussian $\mathcal{N}(\mathbf{x}|\boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)$
 - If $\mathbf{x} \in \{0,1\}^D$, can use D Bernoullis (one for each feature)
 - Amount of training data available ([important](#))
 - If D large and N_k small, it will be difficult to get a good estimate θ_k

E.g., if $p(\mathbf{x}|\theta_k)$ is multivariate Gaussian then assume it to have a diagonal covariance matrix instead of full covariance matrix

Such assumptions greatly reduce the number of parameters to be estimated

In such cases, we may need to [regularize](#) θ_k or make some [simplifying assumptions](#) on $p(\mathbf{x}|\theta_k)$, such as features being conditionally independent given class e.g., $p(\mathbf{x}|\theta_k) = \prod_{d=1}^D p(x_d|\theta_{kd})$ - [naïve Bayes](#)

Especially if the number of features (D) is very large because large value of D means k consists of a large number of parameters (e.g., in the Gaussian case, $\theta_k = (\boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)$, D params for $\boldsymbol{\mu}_k$ and $O(D^2)$ params for $\boldsymbol{\Sigma}_k$). [Can overfit](#)

Generative Classification: At Test Time

- Recall the form of the conditional distribution of the label

$$p(y_* = k | \mathbf{x}_*) = \frac{p(y_* = k) \times p(\mathbf{x}_* | y_* = k)}{p(\mathbf{x}_*)}$$

Class-marginal accounts for the frequency of class k labels in the training data

Class-conditional distribution of inputs accounts for the shape/spread of class k

Probability of \mathbf{x}_* belonging to class k is proportional to the fraction of training inputs from class k times the probability of \mathbf{x}_* under the distribution of inputs from class k

$$\propto p(y_* = k) \times p(\mathbf{x}_* | y_* = k)$$
$$\propto \hat{\pi}_k \times p(\mathbf{x}_* | \hat{\theta}_k)$$

- If we assume the class-marginal to be uniform ($p(y_* = k) = 1/K$) then

$$p(y_* = k | \mathbf{x}_*) \propto p(\mathbf{x}_* | \hat{\theta}_k)$$

- The most likely label is $y_* = \operatorname{argmax}_{k \in \{1, 2, \dots, K\}} p(y_* = k | \mathbf{x}_*)$

Gen. Class. using Gaussian Class-conditionals

- The generative classification model $p(y = k|x) = \frac{p(y=k)p(x|y=k)}{p(x)}$
- Assume each class-conditional $p(x|y = k)$ to be a Gaussian

$$\mathcal{N}(x|\mu_k, \Sigma_k) = \frac{1}{\sqrt{(2\pi)^D |\Sigma_k|}} \exp[-(x - \mu_k)^\top \Sigma_k^{-1} (x - \mu_k)]$$

A benefit of modeling each class by a distribution

Since the Gaussian's covariance models its shape, we can learn the shape of each class

- Class marginal is multinoulli $p(y = k) = \pi_k, \pi_k \in (0,1), \sum_{k=1}^K \pi_k = 1$
- Let's denote the parameters of the model collectively by $\theta = \{\pi_k, \mu_k, \Sigma_k\}_{k=1}^K$
 - Can estimate these using MLE/MAP/Bayesian inference
 - Already saw the MLE solution for $\pi: \pi_k = N_k/N$ (can also do MAP)
 - MLE solution for $\mu_k = \frac{1}{N_k} \sum_{y_n=k} \mathbf{x}_n, \Sigma_k = \frac{1}{N_k} \sum_{y_n=k} (\mathbf{x}_n - \mu_k)(\mathbf{x}_n - \mu_k)^\top$
- If using point est (MLE/MAP) for θ , predictive distribution will be

Can predict the most likely class for the test input \mathbf{x}_* by comparing these probabilities for all values of k

$$p(y_* = k|x_*, \theta) = \frac{\pi_k |\Sigma_k|^{-1/2} \exp\left[-\frac{1}{2}(\mathbf{x}_* - \mu_k)^\top \Sigma_k^{-1} (\mathbf{x}_* - \mu_k)\right]}{\sum_{k=1}^K \pi_k |\Sigma_k|^{-1/2} \exp\left[-\frac{1}{2}(\mathbf{x}_* - \mu_k)^\top \Sigma_k^{-1} (\mathbf{x}_* - \mu_k)\right]}$$

Can also do MAP estimation for μ_k, Σ_k using a Gaussian prior on μ_k and inverse Wishart prior on Σ_k

Exercise: Try to derive this. I will provide a separate note containing the derivation


Note that the exponent has a Mahalanobis distance like term. Also, accounts for the fraction of training examples in class k

Decision Boundary with Gaussian Class-Conditional

- As we saw, the prediction rule when using Gaussian class-conditional

$$p(y = k|x, \theta) = \frac{\pi_k |\Sigma_k|^{-1/2} \exp\left[-\frac{1}{2}(x - \mu_k)^\top \Sigma_k^{-1}(x - \mu_k)\right]}{\sum_{k=1}^K \pi_k |\Sigma_k|^{-1/2} \exp\left[-\frac{1}{2}(x - \mu_k)^\top \Sigma_k^{-1}(x - \mu_k)\right]}$$

- The decision boundary between any pair of classes will be a **quadratic curve**




Reason: For any two classes k and k' at the decision boundary, we will have $p(y = k|x, \theta) = p(y = k'|x, \theta)$. Comparing **their logs** and ignoring terms that don't contain x , can easily see that

$$(x - \mu_k)^\top \Sigma_k^{-1}(x - \mu_k) - (x - \mu_{k'})^\top \Sigma_{k'}^{-1}(x - \mu_{k'}) = 0$$

Decision boundary contains all inputs x that satisfy the above. This is a **quadratic function** of x (this model is sometimes referred to **Quadratic Discriminant Analysis**)

Decision Boundary with Gaussian Class-Conditional


- Assume all classes are modeled using the same covariance matrix $\Sigma_k = \Sigma, \forall k$
- In this case, the decision boundary b/w any pair of classes will be **linear**



Reason: Again using $p(y = k|x, \theta) = p(y = k'|x, \theta)$, comparing their logs and ignoring terms that don't contain \mathbf{x} , we have

$$(\mathbf{x} - \boldsymbol{\mu}_k)^T \boldsymbol{\Sigma}^{-1} (\mathbf{x} - \boldsymbol{\mu}_k) - (\mathbf{x} - \boldsymbol{\mu}_{k'})^T \boldsymbol{\Sigma}^{-1} (\mathbf{x} - \boldsymbol{\mu}_{k'}) = 0$$

Quadratic terms of \mathbf{x} will cancel out; only linear terms will remain; hence decision boundary will be a linear function of \mathbf{x} (**Exercise:** Verify that we can indeed write the decision boundary between this pair of classes as $\mathbf{w}^T \mathbf{x} + b = 0$ where \mathbf{w} and b depend on $\boldsymbol{\mu}_k$, $\boldsymbol{\mu}_{k'}$ and $\boldsymbol{\Sigma}$)



If we assume the covariance matrices of the assumed Gaussian class-conditionals for any pair of classes to be equal, then the learned separation boundary between this pair of classes will be linear; otherwise, quadratic as shown in the figure on left

Generative Models for Regression

- We can even model regression problems using a generative approach
- Note that the output y is not longer discrete (so no notion of a class-conditional)
- However, the basic rule of recovering a **conditional** from **joint** would still apply

$$p(y|\mathbf{x}, \theta) = \frac{p(\mathbf{x}, y|\theta)}{p(\mathbf{x}|\theta)}$$

A benefit of modeling
each class by a
distribution

- Thus we can model the joint distribution $p(\mathbf{x}, y|\theta)$ of features \mathbf{x} and outputs $y \in \mathbb{R}$
 - If features are real-valued the we can model $p(\mathbf{x}, y|\theta)$ using a $(D + 1)$ -dim Gaussian
 - From this $(D + 1)$ -dim Gaussian, we can get $p(y|\mathbf{x}, \theta)$ using Gaussian conditioning formula
 - If joint is Gaussian, any subset of variables (y here), given the rest (\mathbf{x} here) is also a Gaussian!
 - Refer to the Gaussian results from maths refresher slides for the result

References

- **Section 15.3** Kevin Murphy, [Probabilistic Machine Learning: Advanced Topics](#), MIT Press, 2022 (freely available online)
- **Section 1-3** Michael E. Tipping, "[Bayesian inference: An introduction to principles and practice in machine learning.](#)" Summer school on machine learning. Berlin, Heidelberg: Springer Berlin Heidelberg, 2003. 41-62.
- **Chapter 9**, Kevin Murphy, [Probabilistic Machine Learning: An Introduction](#)
- **Lectures 5, 6, 8** Piyush Rai, Probabilistic Machine Learning (CS772A)
- **Lecture 15**, Piyush Rai, Introduction to Machine Learning (CS771A)