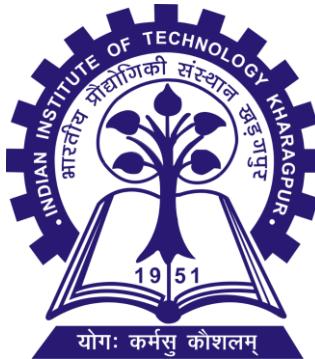


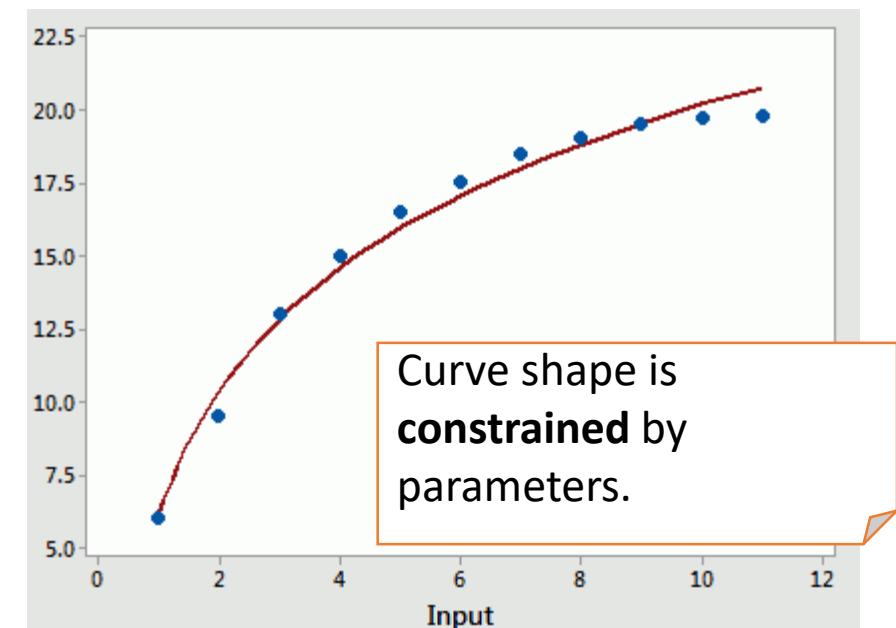
# Gaussian Processes

Debaditya Roy



# Parametric Models

- Discriminative models like logistic regression model  $p(y | \mathbf{x})$  directly
- Functional form is **fixed**:  $p(y|x) = \sigma(\mathbf{w}^\top \mathbf{x} + \mathbf{b})$
- Number of parameters is fixed: Only the weights  $\mathbf{w}$  and bias  $\mathbf{b}$ , no matter how much data we have.
- Limited flexibility: Model can only fit functions within the chosen form (e.g., a linear decision boundary).



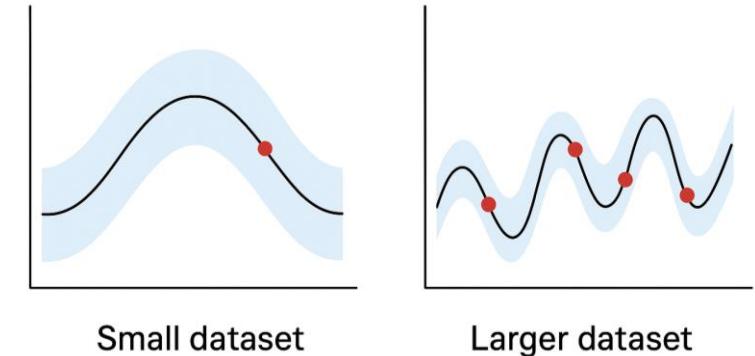
# Non-parametric Models

- Gaussian Processes defines a distribution over functions:

$$f(x) \sim \mathcal{GP}(\mu(x), \kappa(x, x'))$$

- No fixed number of parameters: Predictions are influenced by **all training points**.

- Model complexity grows with data:  
More points  $\rightarrow$  GP adapts its shape.



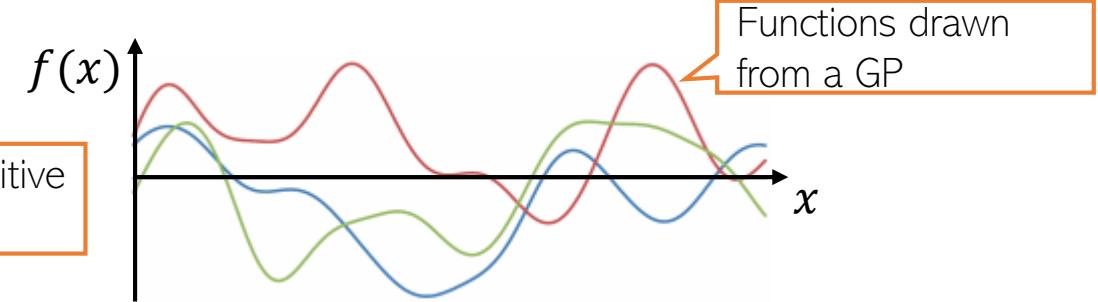
- **Flexible**: Kernel controls smoothness, but GP can model highly non-linear functions.

# Gaussian Process (GP)

- A Gaussian Process (GP) defines a distribution over functions and is denoted as

$$\mathcal{GP}(\mu(\cdot), \kappa(\cdot, \cdot))$$

Mean Function  
 Covariance Function  
 Should be a positive definite function



- Mean function defines what functions drawn from this GP look like on average

$$\mu(x) = \mathbb{E}_{f \sim \mathcal{GP}(\mu, \kappa)} [f(x)]$$

- Covariance/kernel function defines the similarity between a pair of inputs

$$K_{ij} = \kappa(x_i, x_j)$$

- Covariance/kernel function controls the shape of the functions drawn from GP
  - If  $\kappa(x_i, x_j)$  is high then  $f(x_i)$  and  $f(x_j)$  will have similar values

# Covariance/kernel functions

- Kernel functions are popular in learning nonlinear functions (e.g., kernel SVMs)
- Using a kernel corresponds to applying a nonlinear mapping function  $\phi$  on inputs, s.t.,

Kernel function  $\kappa(\dots)$  gives the pairwise similarity between two inputs  $x$  and  $x'$  in the new feature space defined by mapping function  $\phi(\cdot)$

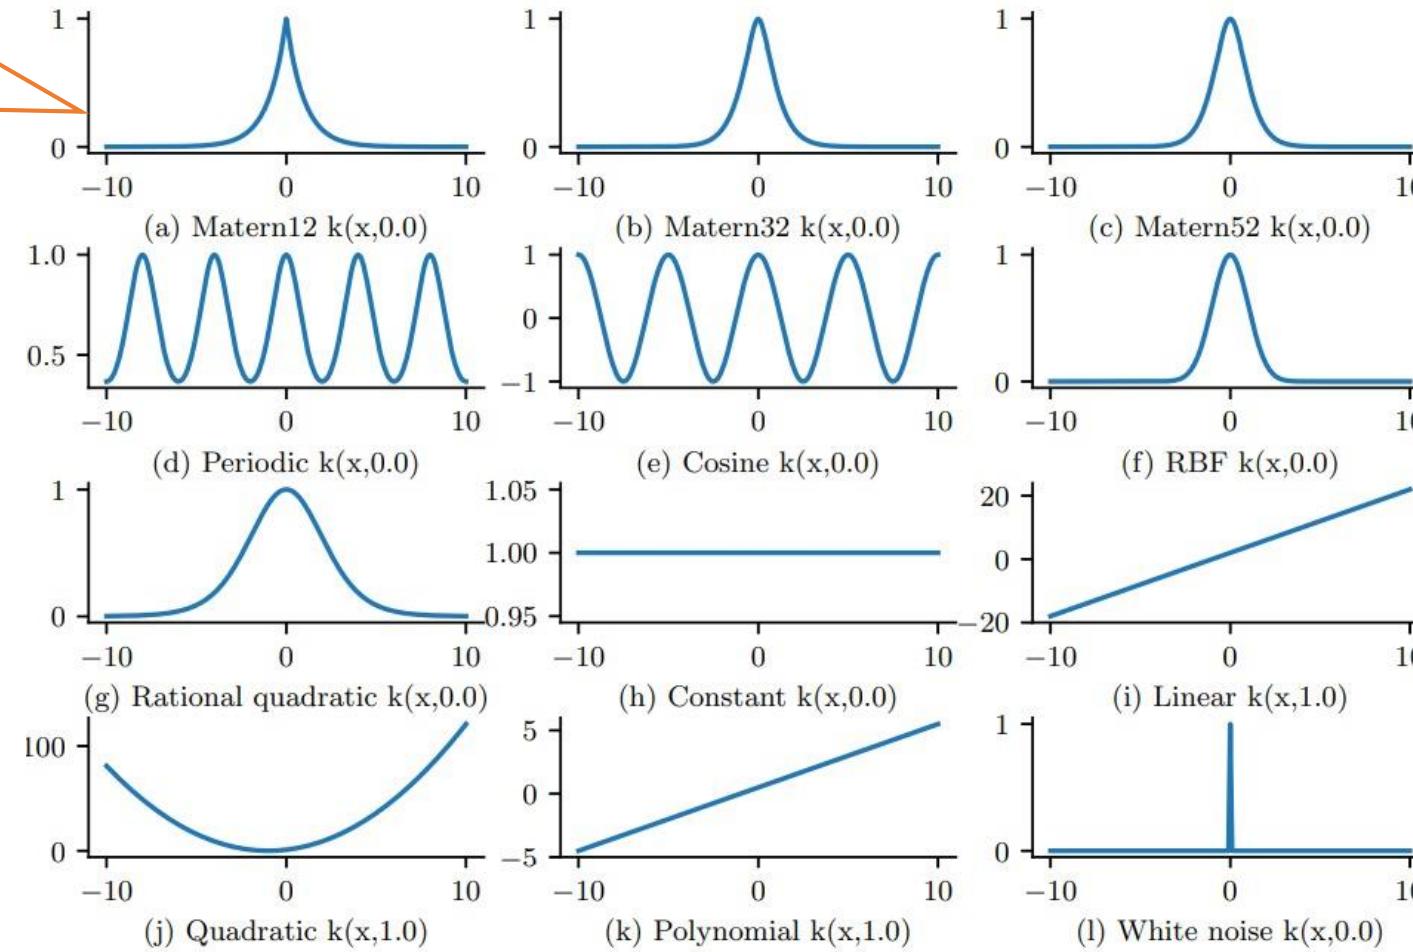
$$\kappa(x, x') = \phi(x)^T \phi(x')$$

- A wide variety of kernel functions exists that suit different types of data
  - Linear kernel, polynomial kernel, Squared exponential (RBF) kernel
  - Automatic Relevance Determination (ARD) kernel
  - Matérn kernel, Periodic kernel, and many others
- We can combine multiple kernels and use them for GP, e.g., possible kernels
  - $\kappa_1 + \kappa_2$ ,  $\kappa_1 \times \kappa_2$ ,  $\alpha\kappa_1 + \kappa_2$ , etc (add, mult., positive scalar mult., or combinations, etc)
- We can learn how to combine multiple kernels and learn their hyperparameters
  - Possible naturally with the Bayesian approach

# Covariance/kernel functions

- Visualization of some kernel functions (how similarity changes with “distance”)

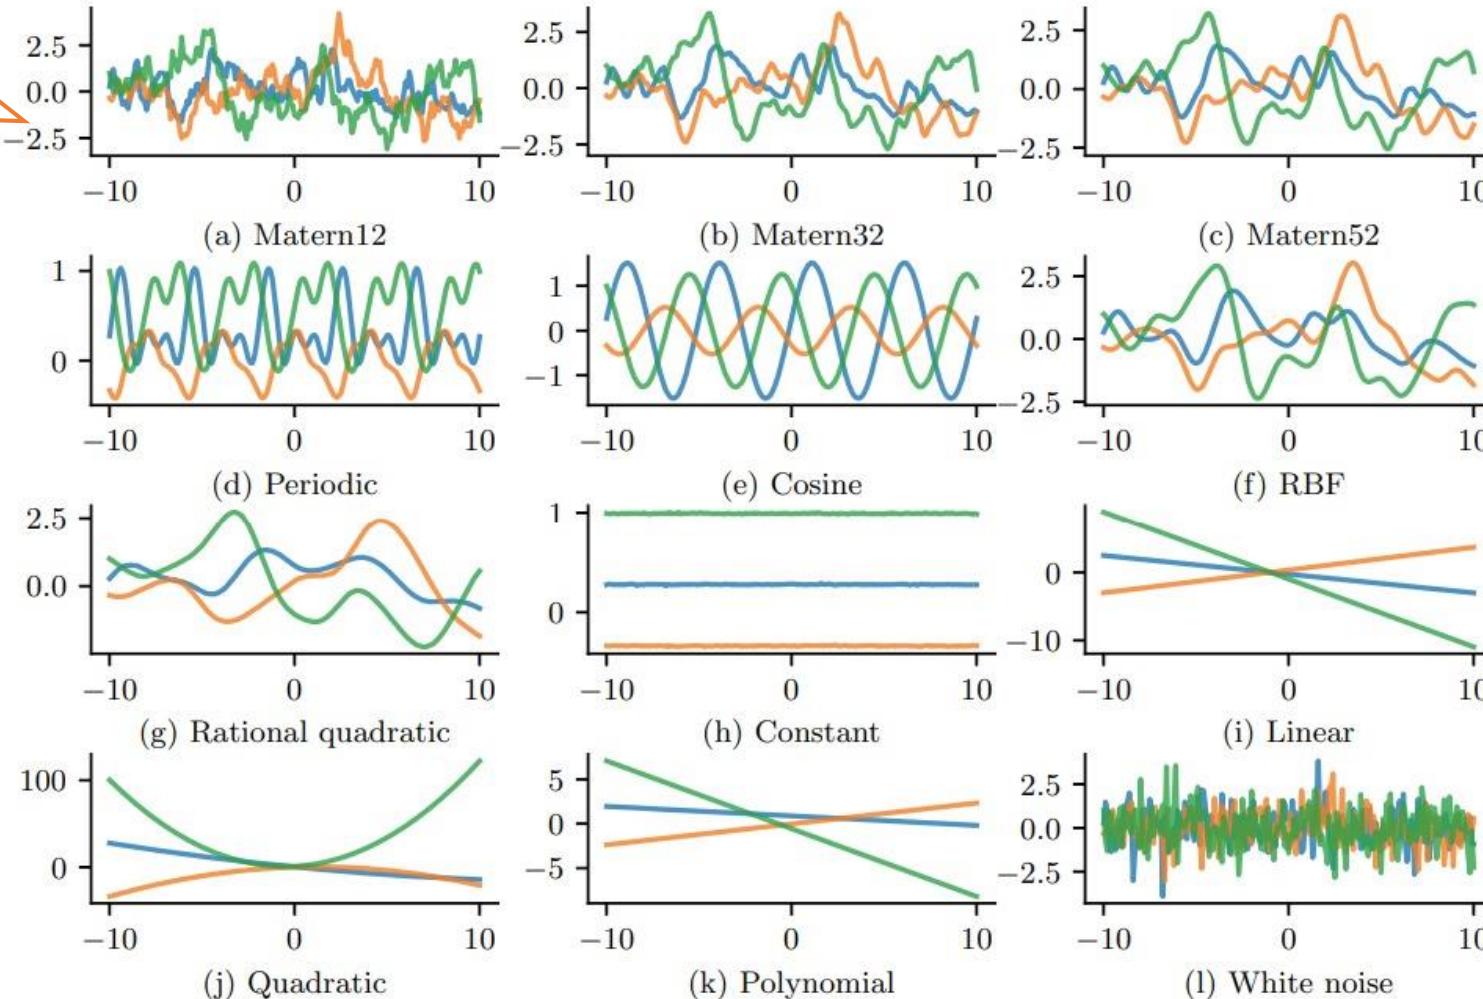
Kernel for all the plots is  $\kappa(x, x')$ , i.e., fixing one input as  $x'$  and varying the other input  $x$  (along the X axis)



# Covariance/kernel functions

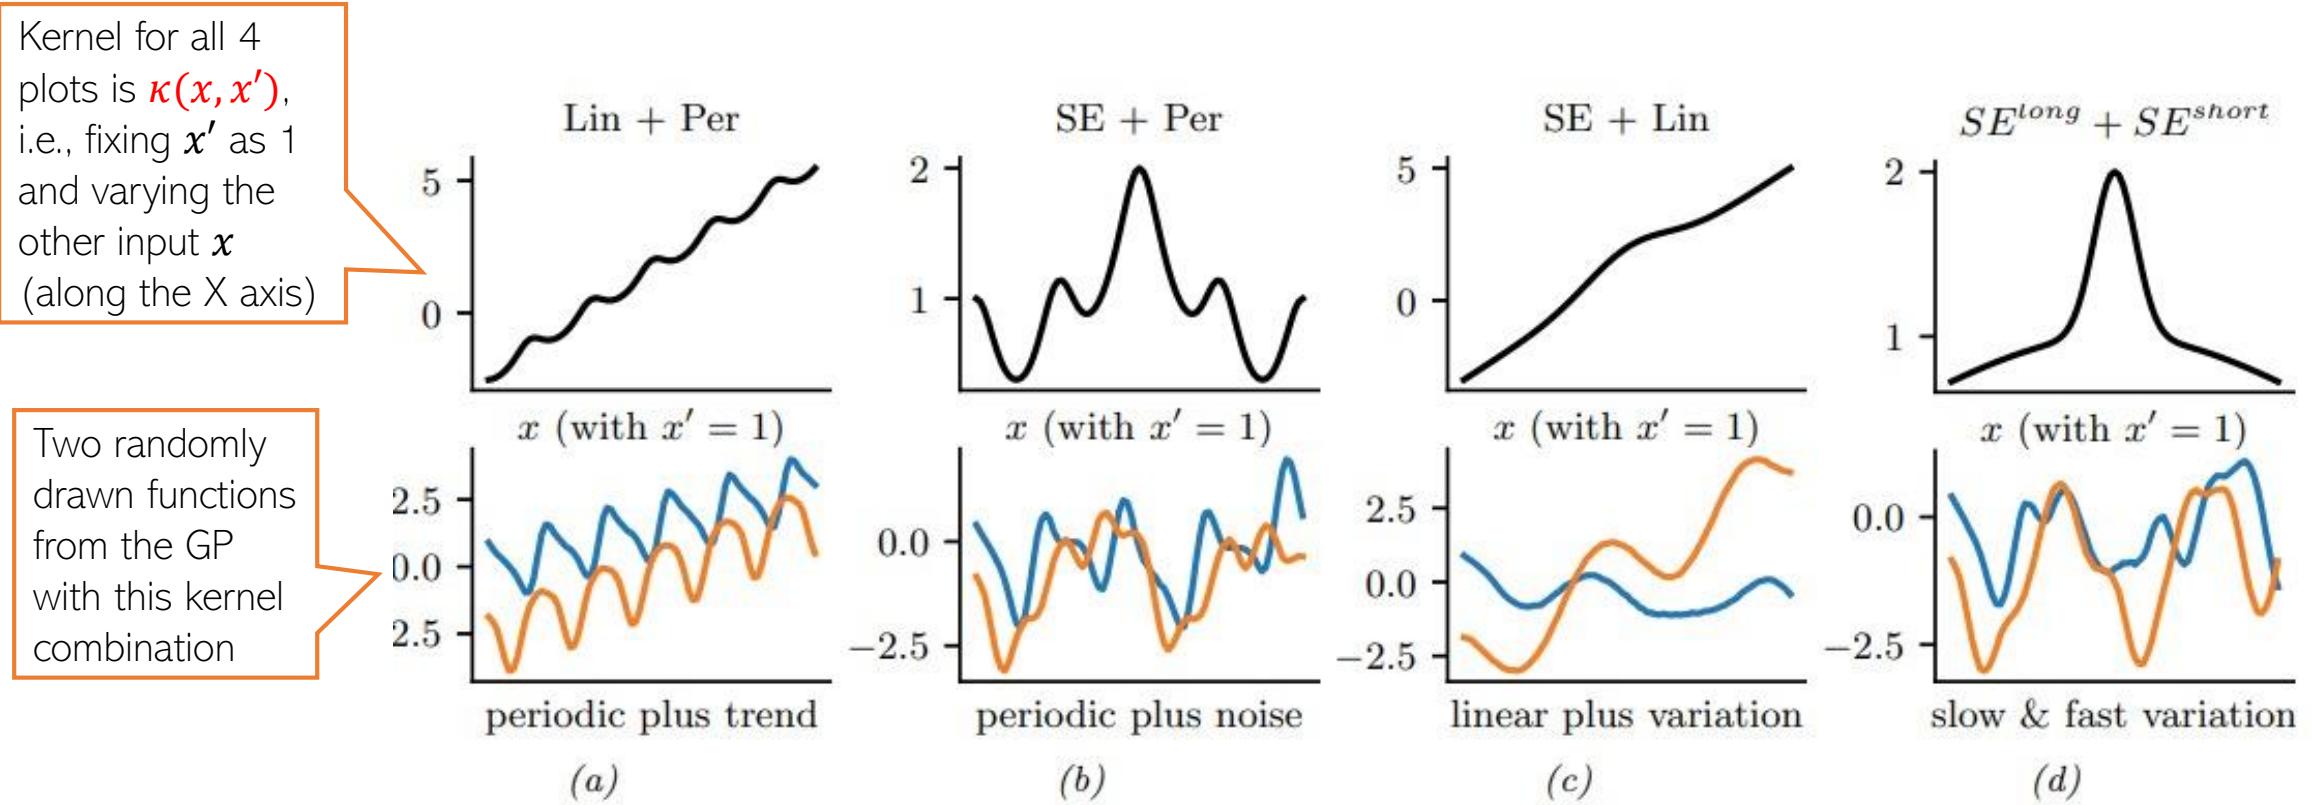
- Examples of functions  $f \sim \text{GP}(\mu, \kappa)$  drawn from a GP using some standard kernels

Each plot shows 3 random functions drawn from the corresponding GP with the specified kernel function



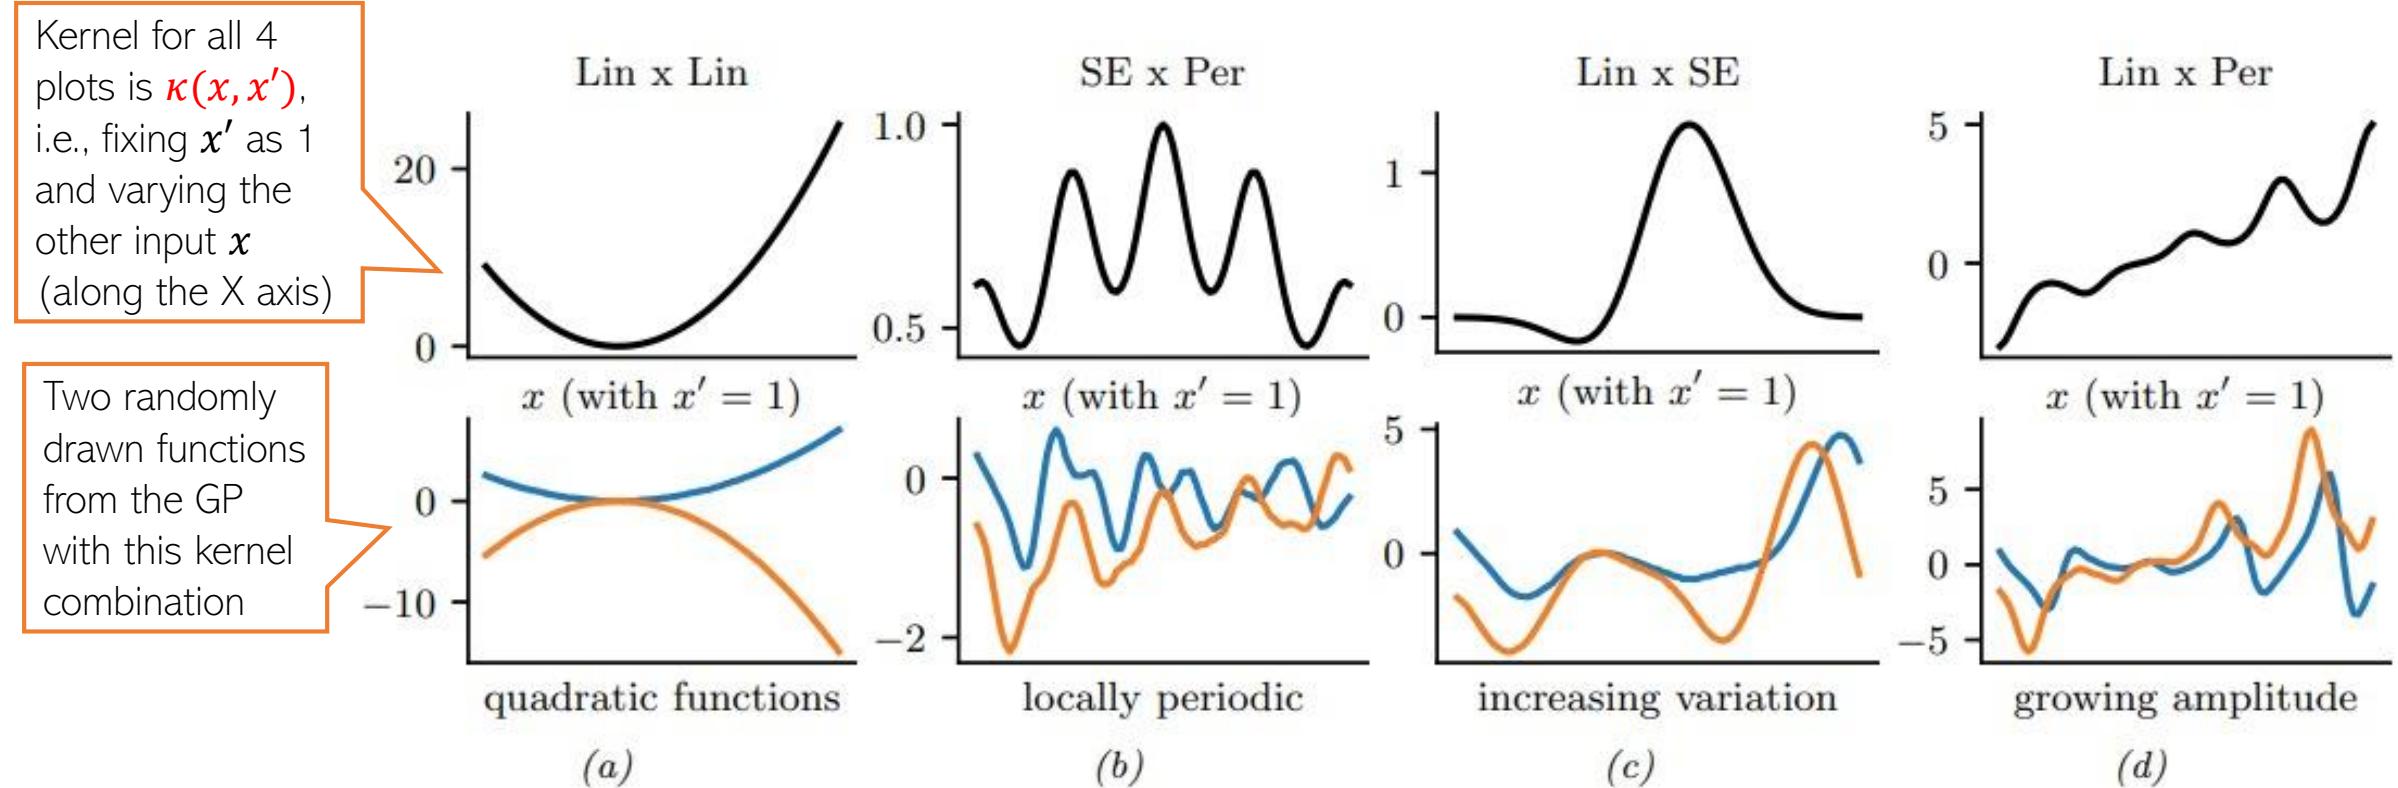
# Combining two (or more) covariance/kernel functions

- Adding two kernels and its effect on the function  $f$  defined by the resulting kernel



# Combining two (or more) covariance/kernel functions

- Multiplying two kernels and its effect on the function  $f$  defined by the resulting kernel



# Gaussian Process: The Predictive Model

- If  $f \sim \mathcal{GP}(\mu, \kappa)$  then  $f$ 's value at any finite set of inputs is jointly Gaussian

$$p\left(\begin{bmatrix} f(x_1) \\ f(x_2) \\ \vdots \\ f(x_N) \end{bmatrix}\right) = \mathcal{N}\left(\begin{bmatrix} \mu(x_1) \\ \mu(x_2) \\ \vdots \\ \mu(x_N) \end{bmatrix}, \begin{bmatrix} \kappa(x_1, x_1) & \dots & \kappa(x_1, x_N) \\ \kappa(x_2, x_1) & \ddots & \kappa(x_2, x_N) \\ \vdots & \ddots & \vdots \\ \kappa(x_N, x_1) & \dots & \kappa(x_N, x_N) \end{bmatrix}\right) \rightarrow p(\mathbf{f}) = \mathcal{N}(\mathbf{f} | \boldsymbol{\mu}, \mathbf{K})$$

N × 1      N × 1      N × N

- Denoting  $f$ 's score for a new test input  $\mathbf{x}_*$  as  $f_* = f(\mathbf{x}_*)$ , we must also have

$$p\left(\begin{bmatrix} \mathbf{f} \\ f_* \end{bmatrix}\right) = \mathcal{N}\left(\begin{bmatrix} \boldsymbol{\mu} \\ \mu_* \end{bmatrix}, \begin{bmatrix} \mathbf{K} & \mathbf{k}_* \\ \mathbf{k}_*^\top & \kappa(x_*, x_*) \end{bmatrix}\right)$$

$\mathbf{k}_* = [\kappa(x_1, x_*), \kappa(x_2, x_*), \dots, \kappa(x_N, x_*)]^\top$   
 $N + 1$  dim Gaussian

- Very useful result: Easy to see that, given the above, the GP predictive distribution

$$p(f_* | \mathbf{f}) = \mathcal{N}(f_* | \color{red}{\mu_* + \mathbf{k}_*^\top \mathbf{K}^{-1}(\mathbf{f} - \boldsymbol{\mu})}, \color{blue}{\kappa(x_*, x_*) - \mathbf{k}_*^\top \mathbf{K}^{-1} \mathbf{k}_*})$$

- Thus score of  $f$  on  $\mathbf{x}_*$  given scores on training inputs has a Gaussian distribution

# Conditional Distribution of a Gaussian

$$\begin{bmatrix} a \\ b \end{bmatrix} \sim \mathcal{N}\left(\begin{bmatrix} \mu_a \\ \mu_b \end{bmatrix}, \begin{bmatrix} \Sigma_{aa} & \Sigma_{ab} \\ \Sigma_{ba} & \Sigma_{bb} \end{bmatrix}\right)$$

$$p(b|a) = \mathcal{N}(\mu_b + \Sigma_{ba}\Sigma_{aa}^{-1}(a - \mu_a), \Sigma_{bb} - \Sigma_{ba}\Sigma_{aa}^{-1}\Sigma_{ab})$$



Show how. Left as exercise

# Sampling in GPs

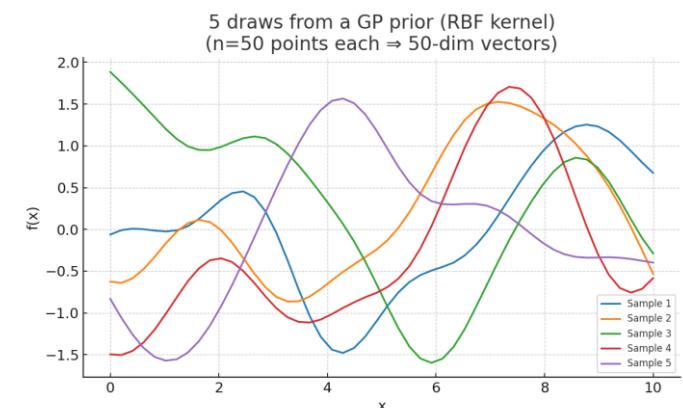
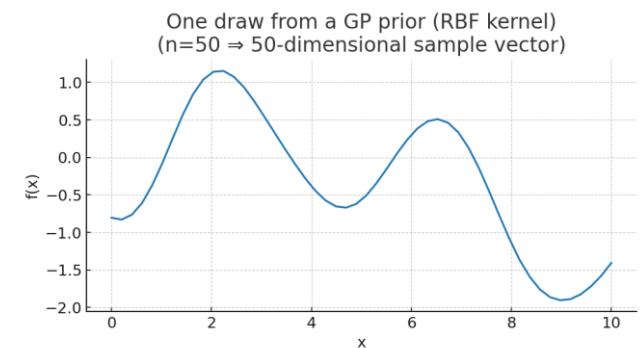
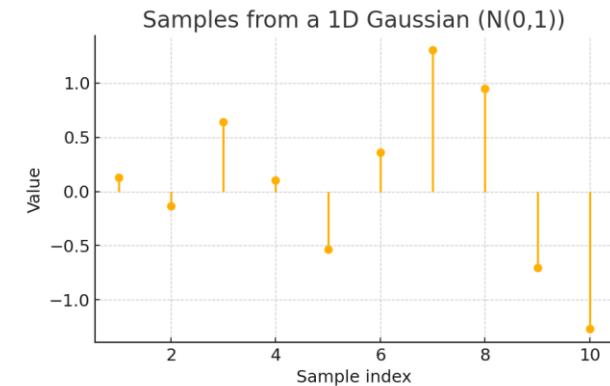
- Pick input points  $X$  (e.g., 50 evenly spaced points in  $[0,10]$ ).
- Compute covariance matrix  $K$  using your kernel  $k(x, x')$

$$k_{SE}(x, x') = \sigma^2 \exp\left(-\frac{(x - x')^2}{2\ell^2}\right)$$

- Draw a sample from the multivariate normal:

$$\begin{aligned} f(X) &\sim \mathcal{N}(0, K) \\ f(X) &= [f(x_1), f(x_2), \dots, f(x_n)]^\top \end{aligned}$$

- Plot the sampled values against  $X$ .
- Each sample corresponds to **one function** from the GP prior.



# Gaussian Process: The Predictive Model

- Assuming the mean function  $\mu(\mathbf{x}) = 0$ , the conditional distribution of score becomes

$$p(f_* | \mathbf{f}) = \mathcal{N}(f_* | \mathbf{k}_*^\top \mathbf{K}^{-1} \mathbf{f}, \kappa(x_*, x_*) - \mathbf{k}_*^\top \mathbf{K}^{-1} \mathbf{k}_*) = \mathcal{N}(f_* | \hat{\mu}, \hat{\sigma}^2)$$

- Note that the predictive mean  $\hat{\mu}$  can be written in the following two equivalent ways

Weighted sum of the scores of the  $N$  training inputs

$$\hat{\mu} = \sum_{i=1}^N \beta_i f_i$$

Like doing a “weighted” nearest neighbors using all the  $N$  training inputs as neighbors with weight  $\beta_i$  given to input  $x_i$  with score  $f_i = f(x_i)$

Weighted sum of kernel based similarities of  $\mathbf{x}_*$  with the  $N$  training inputs

$$\hat{\mu} = \sum_{i=1}^N \alpha_i \kappa(x_i, x_*)$$

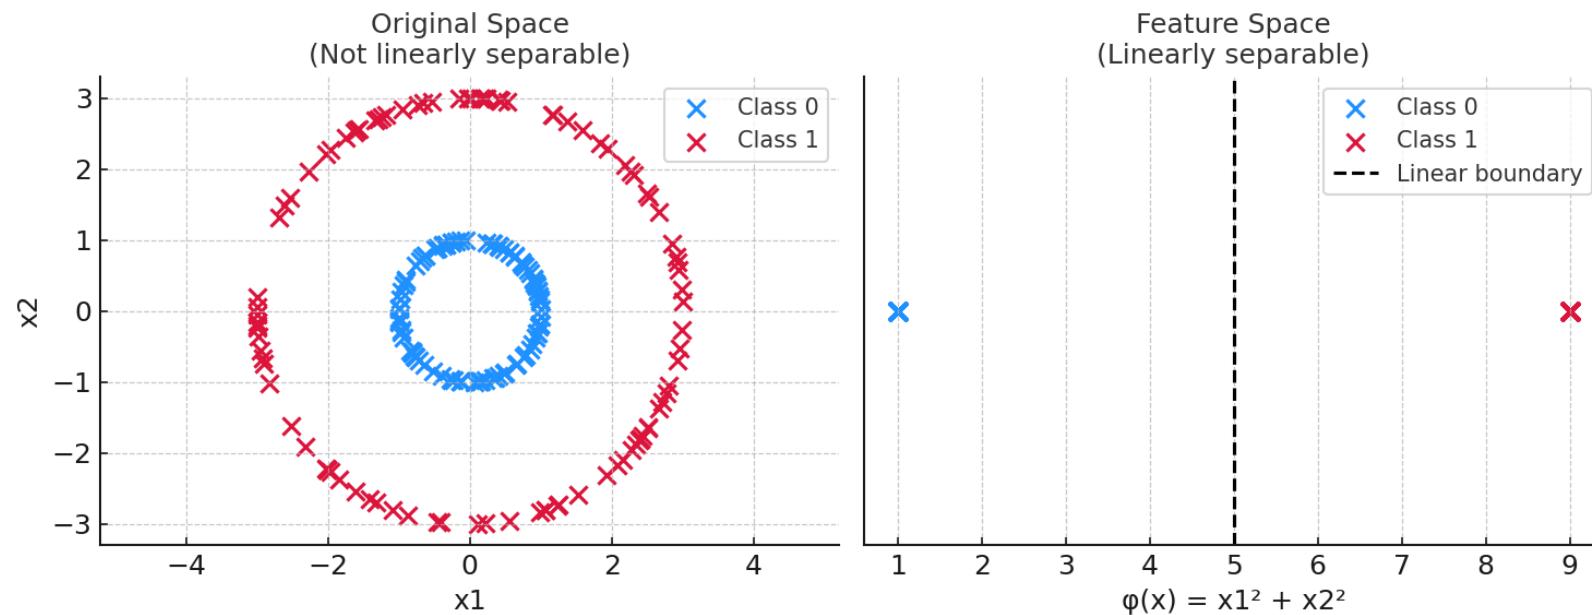
Methods like kernel regression or kernel SVM have their predictions in this form

- Advantage: GP also gives the score's variance  $\hat{\sigma}^2 = \kappa(x_*, x_*) - \mathbf{k}_*^\top \mathbf{K}^{-1} \mathbf{k}_*$

- Thus GP can be viewed as a probabilistic/Bayesian version of kernel methods

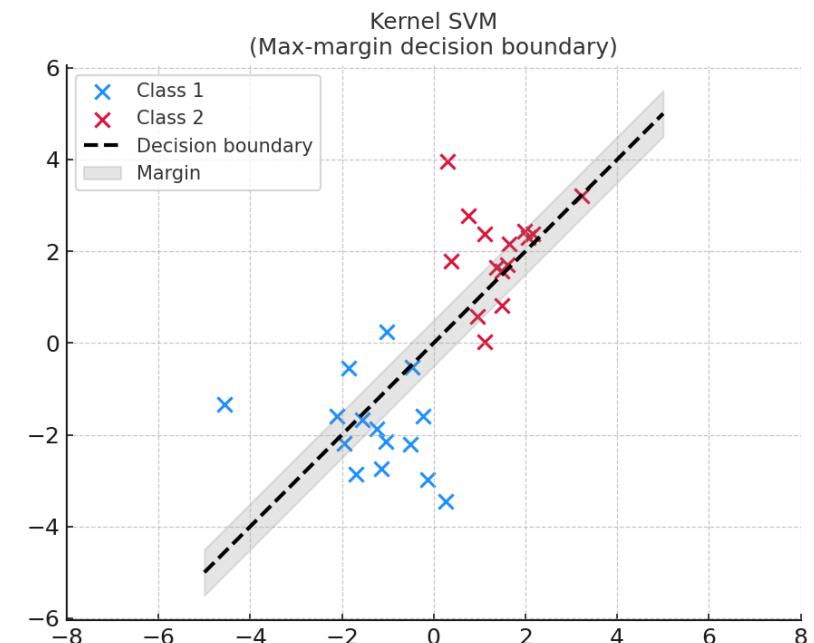
# Detour: Kernel Methods

- Many datasets are not linearly separable
- **Kernel trick:** implicitly map data into higher-dimensional space
- Kernel function:  $K(x, z) = \phi(x) \cdot \phi(z)$   
 $\phi(x) = x_1^2 + x_2^2$



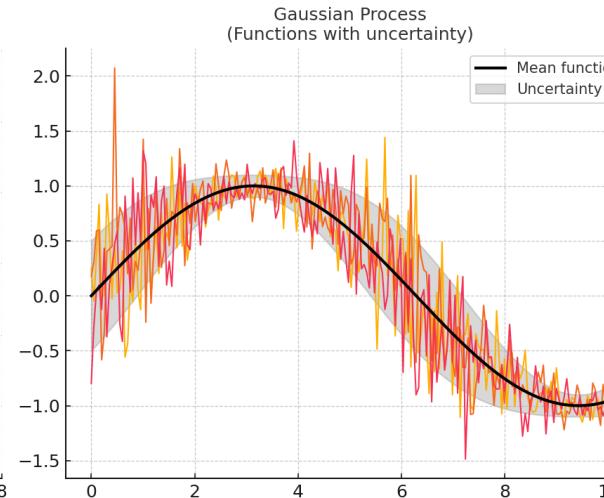
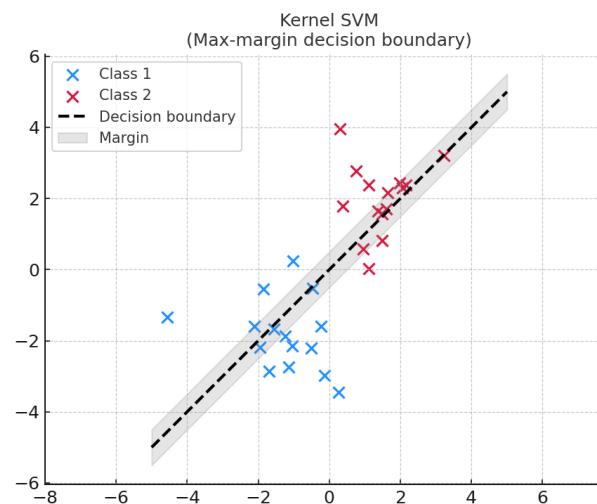
# Kernel SVM

- Support Vector Machine: finds max-margin separating boundary
- Works well with kernels for non-linear decision boundaries
- Key properties:
  - Discriminative model (focuses on boundary)
  - Deterministic outputs (class/margin score)
  - Kernel defines shape of boundary



# Gaussian process

- Each set of points → multivariate normal distribution
- Key properties:
  - Generative (Bayesian) model
  - Probabilistic predictions (mean + uncertainty)
  - Kernel encodes prior assumptions (smoothness, periodicity)
- Difference from SVM:
  - SVM → deterministic boundary
  - GP → distribution of plausible functions



# From GP Scores ( $\mathbf{f}$ ) to Actual Outputs ( $\mathbf{y}$ )

- Assume a supervised learning problem with  $N$  training examples  $(\mathbf{X}, \mathbf{y}) = \{(\mathbf{x}_i, y_i)\}_{i=1}^N$
- Denoting  $f_i = f(x_i)$ , for regression with added noise  $\mathcal{N}(0, \beta^{-1})$

$$y_i = f_i + \epsilon_i$$



$$p(y_i|f_i) = \mathcal{N}(y_i|f_i, \beta^{-1})$$

This GP score  $f_i = f(x_i)$  is the mean of this Gaussian likelihood

The **likelihood** function for all the training outputs (assuming i.i.d.)

$$p(\mathbf{y}|\mathbf{f}) = \mathcal{N}(\mathbf{y}|\mathbf{f}, \beta^{-1}\mathbf{I}_N)$$

For multi-class case with  $C$  classes, we will use a multinoulli with probability vector softmax( $f_i$ ) where  $f_i$  will be a  $C$ -dim vector of logits

- Likewise, for binary classification, **likelihood**  $p(\mathbf{y}|\mathbf{f}) = \prod_{i=1}^N \text{Bernoulli}(y_i|\sigma(f_i))$
- In general, when using GP, the PPD  $p(y_*|\mathbf{y})$  of output  $y_*$  for a new test input  $\mathbf{x}_*$

$$p(y_*|\mathbf{y}) = \int p(y_*|f_*) p(f_*, \mathbf{f}|\mathbf{y}) d\mathbf{f} df_*$$

Averaging over  $f$

The GP posterior

Skipping the training and test inputs from the PPD notation

Likelihood

GP predictive:  
Always a Gaussian

The GP posterior

$$= \int p(y_*|f_*) p(f_*|\mathbf{f}) p(\mathbf{f}|\mathbf{y}) d\mathbf{f} df_*$$

$$p(\mathbf{f}|\mathbf{y}) \propto p(\mathbf{f}) p(\mathbf{y}|\mathbf{f})$$

Likelihood function for training output

The GP prior: Gaussian  $p(\mathbf{f}) = \mathcal{N}(\mathbf{f}|\boldsymbol{\mu}, \mathbf{K})$  or  $p(\mathbf{f}) = \mathcal{N}(\mathbf{f}|0, \mathbf{K})$  is mean function is zero

# GP Prediction with Gaussian Likelihood

- In general, the PPD when using GP is defined as

$$p(y_*|y) = \int p(y_*|f_*) p(f_*|\mathbf{f}) p(\mathbf{f}|y) d\mathbf{f} df_*$$

And don't even have to compute/use the posterior  $p(\mathbf{f}|y)$  (which in this case is a Gaussian by the way 😊 ) to get the PPD

- For Gaussian likelihood (and fixed hyperparams), we don't need to do above integral

- Reason: The marginal likelihood is Gaussian

Gaussian likelihood  
(assuming  $\beta$  is fixed)

$$p(y|\mathbf{f}) = \mathcal{N}(y|\mathbf{f}, \beta^{-1}\mathbf{I}_N)$$

GP prior

Assuming zero  
mean function

$$p(\mathbf{f}) = \mathcal{N}(\mathbf{f} | \mathbf{0}, \mathbf{K})$$

Marginal likelihood  
of training outputs

$$p(y) = \int p(y|\mathbf{f})p(\mathbf{f})d\mathbf{f} = \mathcal{N}(y|\mathbf{0}, \mathbf{K} + \beta^{-1}\mathbf{I}_N) = \mathcal{N}(y|\mathbf{0}, \mathbf{C}_N)$$

Marginal likelihood of  
training and test outputs

$$p\left(\begin{bmatrix} \mathbf{y} \\ y_* \end{bmatrix}\right) = \mathcal{N}\left(\begin{bmatrix} \mathbf{y} \\ y_* \end{bmatrix} \mid \begin{bmatrix} \mathbf{0} \\ 0 \end{bmatrix}, \begin{bmatrix} \mathbf{C}_N & \mathbf{k}_* \\ \mathbf{k}_*^\top & \kappa(x_*, x_*) + \beta^{-1} \end{bmatrix}\right)$$

PPD obtained using  
joint to conditional  
results of Gaussians

$$p(y_*|y) = \mathcal{N}(y_* | \mathbf{k}_*^\top \mathbf{C}_N^{-1} \mathbf{y}, \kappa(x_*, x_*) - \mathbf{k}_*^\top \mathbf{C}_N^{-1} \mathbf{k}_* + \beta^{-1})$$

- $p(y_*|y)$  is almost identical to  $p(f_*|\mathbf{f})$  with  $\mathbf{K}$  replaced by  $\mathbf{C}_N$  + extra  $\beta^{-1}$  noise variance

# Learning Hyperparameters in GP based Models

- Can learn the hyperparameters of the GP prior as well as of the likelihood model
- Assuming  $\mu = 0$ , the hyperparams of GP are cov/kernel function hyperparams

$$\kappa(\mathbf{x}_n, \mathbf{x}_m) = \exp\left(-\frac{||\mathbf{x}_n - \mathbf{x}_m||^2}{\gamma}\right)$$

(RBF kernel)

$$\kappa(\mathbf{x}_n, \mathbf{x}_m) = \exp\left(-\sum_{d=1}^D \frac{(\mathbf{x}_{nd} - \mathbf{x}_{md})^2}{\gamma_d}\right)$$

(ARD kernel)

$$\kappa(\mathbf{x}_n, \mathbf{x}_m) = \kappa_{\theta_1}(\mathbf{x}_n, \mathbf{x}_m) + \kappa_{\theta_2}(\mathbf{x}_n, \mathbf{x}_m) + \dots + \kappa_{\theta_M}(\mathbf{x}_n, \mathbf{x}_m) \quad (\text{flexible composition of multiple kernels})$$

Can help in feature selection (irrelevant features will tend to have very large  $\gamma_d$ )Different RBF kernel bandwidth  $\gamma_d$  for each feature

Ability to learn kernel hyperparams (without cross-validation) is another very appealing property of GP



- MLE-II is a popular choice for learning these hyperparams (otherwise MCMC, VI, etc)
- Denoting the covariance/kernel matrix as  $\mathbf{K}_\theta$ , for Gaussian likelihood case, the marg-lik

$$p(\mathbf{y}|\theta, \beta^{-1}) = \mathcal{N}(\mathbf{y}|\mathbf{0}, \mathbf{K}_\theta + \beta^{-1}\mathbf{I}_N)$$

- This can be maximized to learn  $\theta$  and  $\beta$
- For non-Gaussian likelihoods, the marg-lik itself will need to be approximated

# Weight Space View vs Function Space View

- GPs are defined w.r.t. a **function space** that models input-output relationship
- In contrast, we have seen models that are defined w.r.t. a **weight space**, e.g.,

$$p(\mathbf{y}|\mathbf{X}, \mathbf{w}) = \mathcal{N}(\mathbf{y}|\mathbf{X}\mathbf{w}, \beta^{-1}\mathbf{I}_N)$$

Likelihood

$$p(\mathbf{w}) = \mathcal{N}(\mathbf{w}|\boldsymbol{\mu}_0, \boldsymbol{\Sigma}_0)$$

Prior over weight vector

$$p(\mathbf{y}|\mathbf{X}) = \int p(\mathbf{y}|\mathbf{X}, \mathbf{w})p(\mathbf{w})d\mathbf{w} = \mathcal{N}(\mathbf{y}|\mathbf{X}\boldsymbol{\mu}_0, \beta^{-1}\mathbf{I}_N + \mathbf{X}\boldsymbol{\Sigma}_0\mathbf{X}^\top)$$

Marginal likelihood  
after integrating  
out the weights

$$p(\mathbf{y}|\mathbf{X}) = \mathcal{N}(\mathbf{y}|\mathbf{0}, \beta^{-1}\mathbf{I}_N + \mathbf{X}\mathbf{X}^\top)$$

Marginal likelihood assuming  $\boldsymbol{\mu}_0 = \mathbf{0}$  and  $\boldsymbol{\Sigma}_0 = \mathbf{I}$ 

$$p(\mathbf{y}|\mathbf{X}) = \mathcal{N}(\mathbf{y}|\mathbf{0}, \mathbf{X}\mathbf{X}^\top)$$

Assuming noise-free likelihood

Linear-Gaussian transformation

```

Write the model equation as
y = Xw + e, e ~ N(0, I_N), w ~ N(\mu_0, \Sigma_0)

with w independent of e
and the transformation of e leaves it unchanged with
+ Mean: E[e] = 0
+ Covariance: Cov[e] = E[e e^\top] = E[e e] = E[e^2] - E[e]^2
+ Noise: Cov[y] = E[(y - E[y])(y - E[y])^\top] = E[(y - E[y])^2]

```

- Thus the **joint marginal** of the  $N$  responses  $y_1, y_2, \dots, y_N$  is a **multivariate Gaussian**

This equivalence also shows that Bayesian linear regression is a special case of GP with linear kernel

$$p\left(\begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_N \end{bmatrix}\right) = \mathcal{N}\left(\begin{bmatrix} 0 \\ 0 \\ \vdots \\ 0 \end{bmatrix}, \begin{bmatrix} x_1^\top x_1 & \dots & x_1^\top x_N \\ x_2^\top x_1 & \dots & x_2^\top x_N \\ \vdots & \ddots & \vdots \\ x_N^\top x_1 & \dots & x_N^\top x_N \end{bmatrix}\right)$$

Same as a GP  $f(x_i) = y_i$ ,  $\mu(x) = \mathbf{0}$  and linear covariance/kernel function  $\kappa(x_i, x_j) = x_i^\top x_j$

- Thus GPs can be seen as bypassing the weight space and directly defining the model using a marginal likelihood via a function space defined by the GP

# Linear-Gaussian transformation

Write the model explicitly as

$$\mathbf{y} = X\mathbf{w} + \boldsymbol{\varepsilon}, \quad \boldsymbol{\varepsilon} \sim \mathcal{N}(\mathbf{0}, \beta^{-1}I_N), \quad \mathbf{w} \sim \mathcal{N}(\boldsymbol{\mu}_0, \Sigma_0),$$

with  $\mathbf{w}$  independent of  $\boldsymbol{\varepsilon}$ .

An affine transformation of Gaussians is Gaussian, so  $\mathbf{y}$  is Gaussian with

- Mean:

$$\mathbb{E}[\mathbf{y}] = X \mathbb{E}[\mathbf{w}] + \mathbb{E}[\boldsymbol{\varepsilon}] = X\boldsymbol{\mu}_0 + \mathbf{0}.$$

- Covariance:

$$\text{Cov}[\mathbf{y}] = X \text{Cov}[\mathbf{w}] X^\top + \text{Cov}[\boldsymbol{\varepsilon}] = X\Sigma_0X^\top + \beta^{-1}I_N.$$

Hence

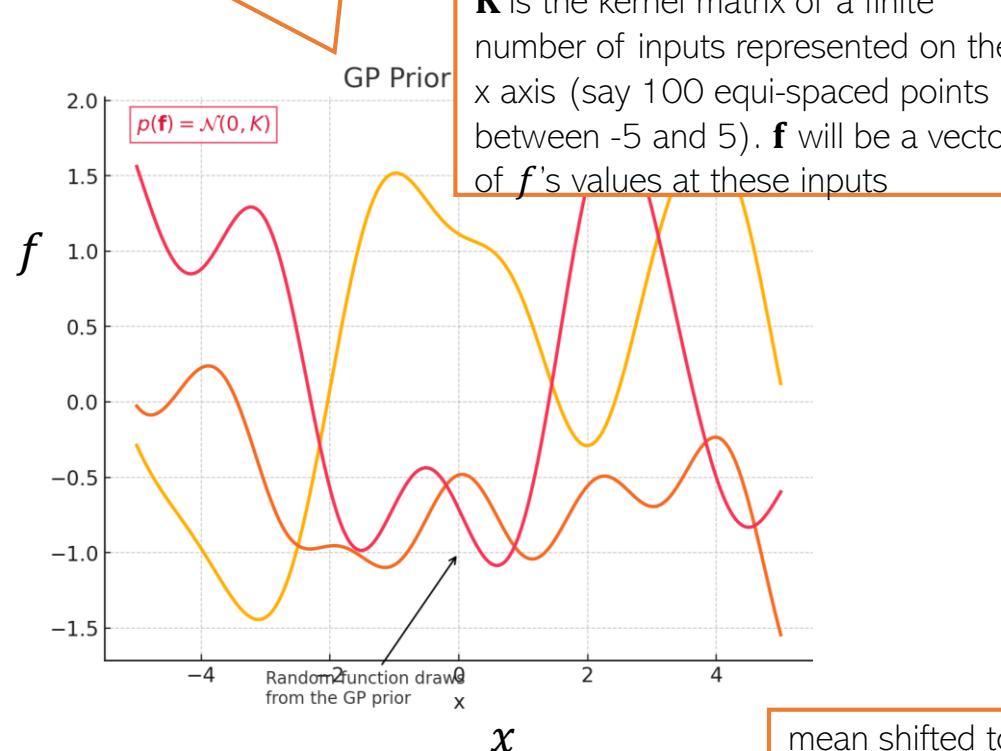
$$\boxed{\mathbf{y} \sim \mathcal{N}\left(X\boldsymbol{\mu}_0, X\Sigma_0X^\top + \beta^{-1}I_N\right)}.$$

# GP: A Visualization

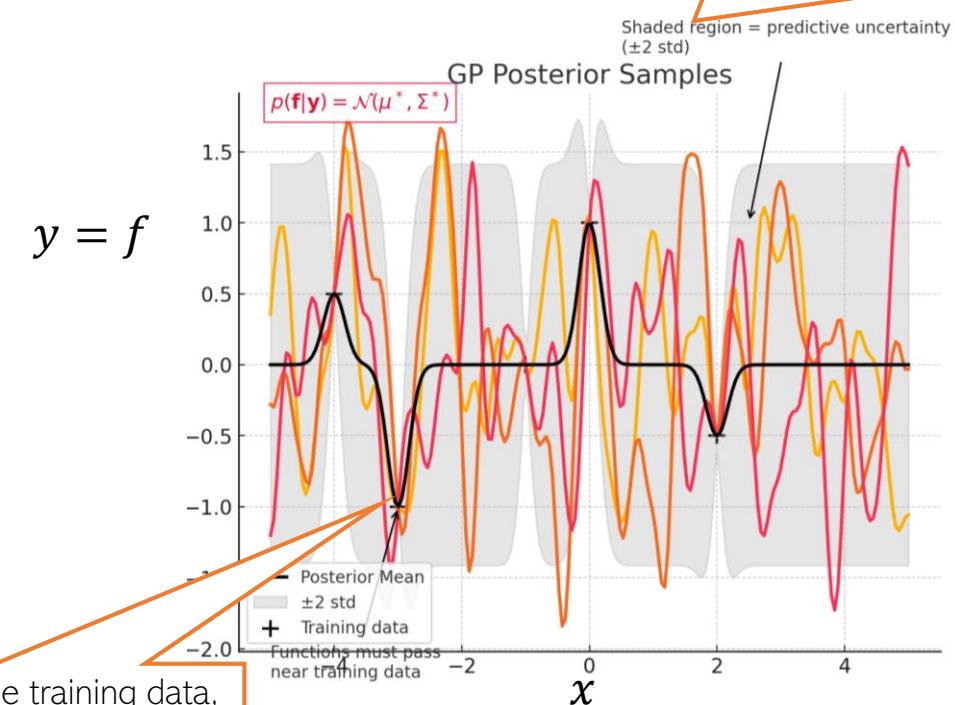
- Assume zero mean function and a squared exponential kernel

$$k_{SE}(x, x') = \sigma^2 \exp\left(-\frac{(x-x')^2}{2\ell^2}\right)$$

Each curve below is obtained by drawing a random  $\mathbf{f}$  from the GP prior  $p(\mathbf{f}) = \mathcal{N}(\mathbf{0}, \mathbf{K})$  and plotting it.



Each curve below is obtained by drawing random  $\mathbf{f}$ 's from the GP posterior  $p(\mathbf{f}|\mathbf{y})$  which is also a Gaussian (The + symbols denote the training data and we assume noiseless outputs, i.e.,  $y_i = f_i$ ).



# Scalability of GPs

- Computational costs in some steps of GP models scale in the size of training data
- For example, prediction cost is  $O(N)$

$$p(y_*|y) = \mathcal{N}(y_*|\hat{\mu}, \hat{\sigma}^2) \quad \hat{\mu} = \mathbf{k}_*^\top \mathbf{C}_N^{-1} \mathbf{y} \quad \hat{\sigma}^2 = \kappa(x_*, x_*) - \mathbf{k}_*^\top \mathbf{C}_N^{-1} \mathbf{k}_* + \beta^{-1}$$

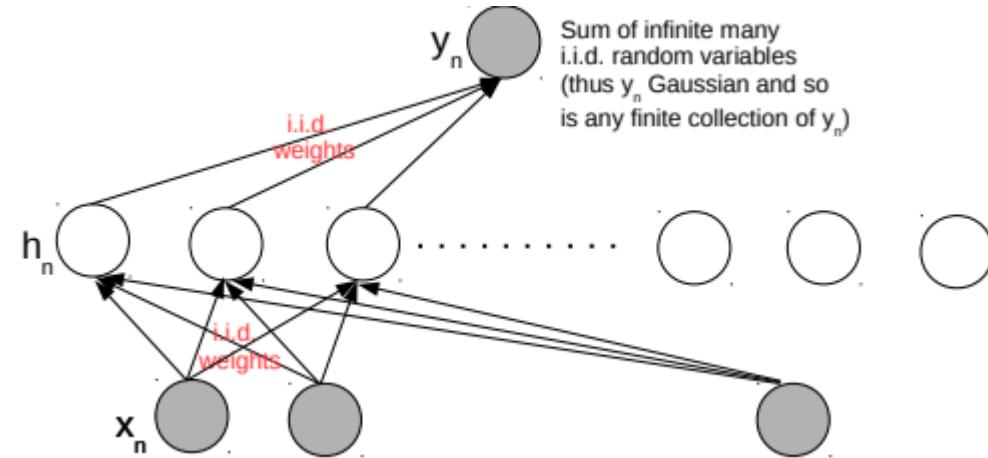
$O(N)$  cost assuming  $\mathbf{C}_N$   
is already inverted

- GP models often require matrix inversions (e.g., in marginal-likelihood computation when estimating hyperparameters) – takes  $O(N^3)$
- Storage also requires  $O(N^2)$  since need to store the covariance matrix
- A lot of work on speeding up GPs<sup>1</sup>. Some prominent approaches include
  - Inducing Point Methods (condition predictions only on a small set of “learnable” points)
  - Divide-and-Conquer (learn GP on small subsets of data and aggregate predictions)
  - Kernel approximations
- Note that nearest neighbor methods and kernel methods also face similar issues
  - Many tricks to speed up kernel methods can be used for speeding up GPs too

$M \ll N$  pseudo-inputs  
and pseudo-outputs

# Neural Networks and Gaussian Process

- An infinitely-wide single hidden layer NN with i.i.d. priors on weights = GP
- Shown formally by (Neal<sup>2</sup>, 1994). Based on applying the central limit theorem



- This equivalence is useful for several reasons
  - Can use a GP instead of an **infinitely wide** Bayesian NN (which is impractical anyway)
  - With GPs, inference is easy (at least for regression and with known hyperparams)
  - A proof that GPs can also learn any function (just like infinitely wide neural nets - Hornik's theorem)
- Connection generalized to infinitely wide multiple hidden layer NN (Lee et al<sup>3</sup>, 2018)

<sup>2</sup>Priors for infinite networks, Tech Report, 1994

<sup>3</sup>Deep Neural Networks as Gaussian Processes (ICLR 2018)

# Infinite Width NNs converge to GPs

## 1. Neural network setup

Consider a simple feedforward neural network (1 hidden layer, scalar output):

$$f(x) = \frac{1}{\sqrt{H}} \sum_{i=1}^H v_i \phi(w_i^\top x + b_i),$$

- $H$ : number of hidden units.
- $w_i, b_i$ : weights and biases of hidden layer.
- $v_i$ : weights to the output.
- $\phi(\cdot)$ : activation function.
- Assume all weights are i.i.d. Gaussian:

$$w_i \sim \mathcal{N}(0, \sigma_w^2 I), \quad b_i \sim \mathcal{N}(0, \sigma_b^2), \quad v_i \sim \mathcal{N}(0, \sigma_v^2).$$

# Infinite Width NNs converge to GPs

## 2. Distribution over outputs as $H \rightarrow \infty$

- For fixed input  $x$ ,  $f(x)$  is a sum of many i.i.d. random terms.
- By the Central Limit Theorem (CLT), as  $H \rightarrow \infty$ ,  $f(x)$  becomes Gaussian distributed.

Now, for multiple inputs  $(x_1, \dots, x_N)$ , the joint distribution of  $(f(x_1), \dots, f(x_N))$  is **multivariate Gaussian** because each pair of outputs shares weights in common.

Thus, in the infinite-width limit:

$$f(x) \sim \text{GP}(0, k(x, x')),$$

where the kernel  $k$  comes from the covariance structure of the NN.

# Infinite Width NNs converge to GPs

## 3. The induced GP kernel

The kernel arises from the expectation over the random weights:

$$k(x, x') = \mathbb{E}[f(x)f(x')].$$

For the 1-hidden-layer network:

$$k(x, x') = \sigma_v^2 \mathbb{E}_{w,b} [\phi(w^\top x + b) \phi(w^\top x' + b)].$$

- Different activations  $\phi$  give different kernels:
  - Linear  $\phi(z) = z \rightarrow$  linear kernel.
  - ReLU  $\rightarrow$  ArcCos kernel.
  - Tanh  $\rightarrow$  "Neal kernel" (from Neal, 1996).

This is sometimes called the **Neural Network Gaussian Process (NNGP)** kernel.

# GP: Some other comments

- GPs can be thought of as Bayesian analogues of kernel methods
- Can get estimate of the uncertainty in the function and its predictions
- Can learn the kernel (by learning the hyperparameters of the kernels)
- In some ways, GPs and (Bayesian/ensembles of) deep neural nets have same goals
  - These methods are also very related (though appear different based on their formulation)
  - Several recent papers have investigated these connections
- GP can be a nice alternative to (Bayesian/ensembles of) deep neural networks
  - GP may be preferable if we don't have that much training data (deep networks requires lots of data to train well)
  - When we have lots of training data, training and test speed may be an issue for GP (but faster versions exist)
- Not limited to supervised learning problems
  - $f$  could even define a mapping of low-dim latent variable  $\mathbf{z}_n$  to an observation  $\mathbf{x}_n$

$$\mathbf{x}_n = f(\mathbf{z}_n) + \text{"noise"}$$

GP latent variable model for dimensionality reduction  
(like a kernel version of probabilistic PCA)

# References

- Chapter 18, Kevin Murphy, [Probabilistic Machine Learning: Advanced Topics](#), MIT Press, 2022 (freely available online)

# GP packages

- Many mature implementations of GP exist. You may check out
  - GPyTorch (PyTorch), GPFlow (Tensorflow)
  - sklearn (Python with some basic GP implementations)
  - GPML (MATLAB), GPsuff (MATLAB/Octave)
  - Many others such as Stan, GPJax
- A comparison of the various packages:  
[https://en.wikipedia.org/wiki/Comparison\\_of\\_Gaussian\\_process\\_software](https://en.wikipedia.org/wiki/Comparison_of_Gaussian_process_software)