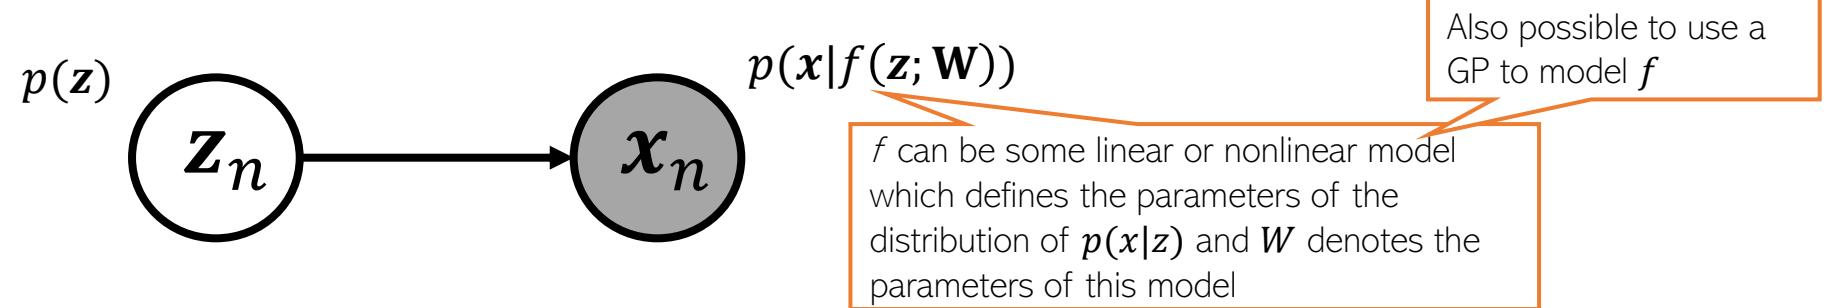


VAE and GAN

Latent Variable Models for Generation Tasks

- Assume a K -dim latent variable \mathbf{z}_n is transformed to generate to D -dim observation \mathbf{x}_n

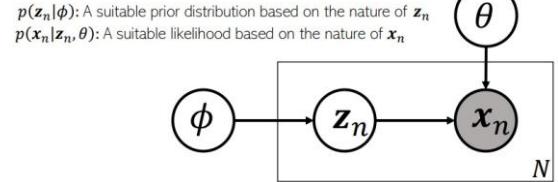


- It is common to use a Gaussian prior for \mathbf{z}_n (though other priors can be used)
- If we use a neural net or GP, such models can generate very high-quality data
 - Take the trained network, generate a random \mathbf{z} from prior, pass it through the model to generate \mathbf{x}



Some sample images generated by Vector Quantized Variational Auto-Encoder (VQ-VAE), a state-of-the-art latent variable model for generation

Factor Analysis and Probabilistic PCA



- FA and PPCA assume f to be a linear model
- In FA/PPCA, latent variables $\mathbf{z}_n \in \mathbb{R}^K$ typically assumed to have a Gaussian prior
 - If we want sparse latent variable, can use Laplace prior on \mathbf{z}_n
 - More complex extensions of FA/PPCA use a mixture of Gaussians prior on \mathbf{z}_n
- Assumption: Observations $\mathbf{x}_n \in \mathbb{R}^D$ typically assumed to have a Gaussian likelihood
 - Other likelihood models (e.g., exp-family) can also be used if data not real-valued
- Relationship between \mathbf{z}_n and \mathbf{x}_n modeled by a noisy linear mapping

$$\mathbf{x}_n = \mathbf{W}\mathbf{z}_n + \epsilon_n = \sum_{k=1}^K \mathbf{w}_k z_{nk} + \epsilon_n$$

Zero-mean and diagonal or spherical Gaussian noise

Linear combination of the columns of \mathbf{W}

$$p(\mathbf{z}_n) = \mathcal{N}(\mathbf{z}_n | \mathbf{0}, \mathbf{I})$$

$$p(\mathbf{x}_n | \mathbf{z}_n) = \mathcal{N}(\mathbf{x}_n | \mathbf{W}\mathbf{z}_n, \Psi)$$

Diagonal for FA, spherical for PPCA

- Linear Gaussian Model. \mathbf{W} , \mathbf{z}_n 's, and Ψ can be learned (e.g, using EM, VI, MCMC)

Some Variants of FA/PPCA

■ Gamma-Poisson latent factor model

Popular for modeling count-valued data (in text analysis, recommender systems, etc)

Non-negative priors often give a nice interpretability to such latent variable models (will see some more examples of such models shortly)

- Assumes K -dim non-negative latent variable \mathbf{z}_n and D -dim count-valued observations \mathbf{x}_n
- An example: Each \mathbf{x}_n is the word-count vector representing a document

$$p(\mathbf{z}_n) = \prod_{k=1}^K \text{Gamma}(z_{nk} | a_k, b_k)$$

$$p(\mathbf{x}_n | \mathbf{z}_n) = \prod_{d=1}^D \text{Poisson}(x_{nd} | f(\mathbf{w}_d, \mathbf{z}_n))$$

This is the rate of the Poisson. It should be non-negative, $\exp(\mathbf{w}_d^\top \mathbf{z}_n)$, or simply $\mathbf{w}_d^\top \mathbf{z}_n$ if \mathbf{w}_d is also non-negative (e.g., using a gamma/Dirichlet prior on it)

- This can be thought of as a probabilistic non-negative matrix factorization model

■ Dirichlet-Multinomial/Multinoulli PCA

- Assumes K -dim non-negative latent variable \mathbf{z}_n and D categorical obs $\mathbf{x}_n = \{\mathbf{x}_{nd}\}_{d=1}^D$
- An example: Each \mathbf{x}_n is a document with D words in it (each word is a categorical value)

Also sums to 1

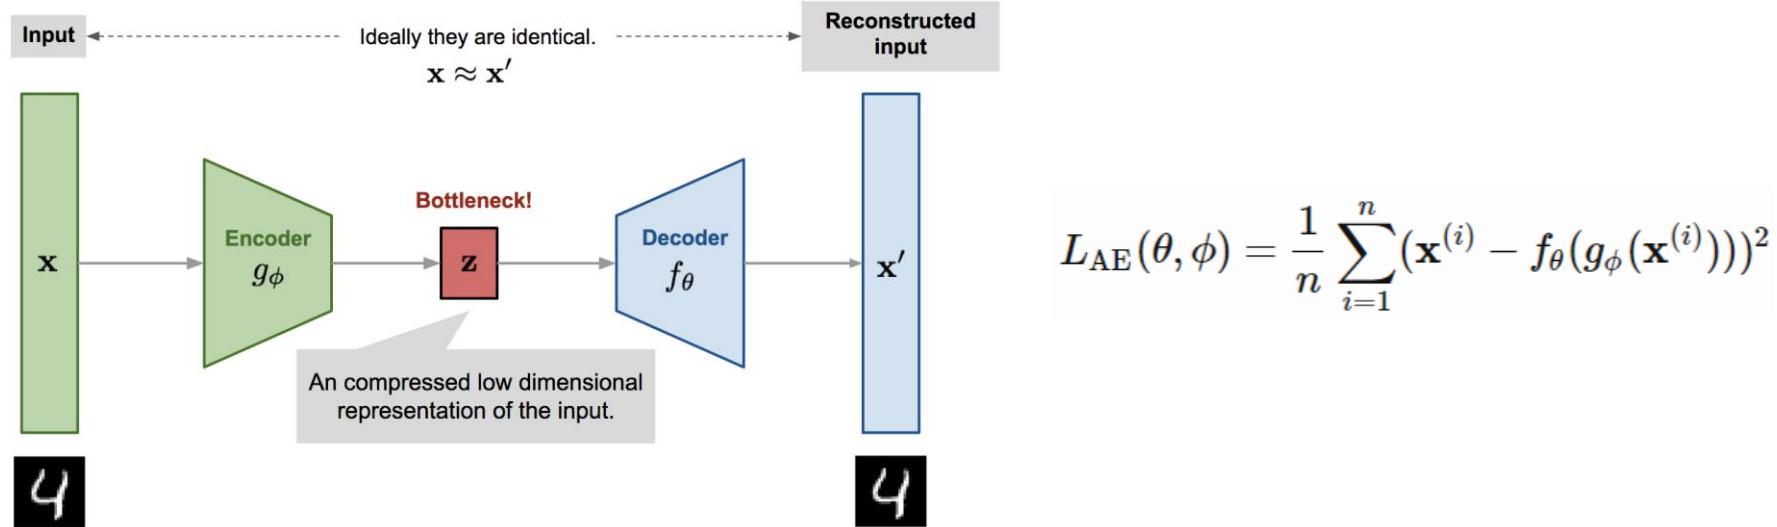
$$p(\mathbf{z}_n) = \text{Dirichlet}(\mathbf{z}_n | \boldsymbol{\alpha})$$

$$p(\mathbf{x}_n | \mathbf{z}_n) = \prod_{d=1}^D \text{Multinoulli}(x_{nd} | f(\mathbf{w}_d, \mathbf{z}_n))$$

This should give the probability vector of the multinoulli over x_{nd} . It should be non-negative and should sum to 1

A Deep Generative Model: Variational Auto-encoder (VAE)

- VAE* is a probabilistic extension of autoencoders (AE). An AE is shown below



- The basic difference is that VAE assumes a prior $p(\mathbf{z})$ on the latent code \mathbf{z}
 - This enables it to not just compress the data but also generate synthetic data
 - How: Sample \mathbf{z} from a prior and pass it through the decoder
- Thus VAE can learn good latent representation + generate novel synthetic data
- The name has “Variational” in it since it is learned using VI principles

Variational Autoencoder (VAE)

- VAE has three main components
 - A prior $p_\theta(\mathbf{z})$ over latent codes
 - A probabilistic decoder/generator $p_\theta(\mathbf{x}|\mathbf{z})$, modeled by a deep neural net
 - A posterior or probabilistic encoder $p_\theta(\mathbf{z}|\mathbf{x})$ approx. by an “inference network” $q_\phi(\mathbf{z}|\mathbf{x})$

- VAE is learned by maximizing the ELBO

ELBO for a single data point

$$\begin{aligned}\mathcal{L}(\theta, \phi | \mathbf{x}) &= \mathbb{E}_{q_\phi(\mathbf{z}|\mathbf{x})} [\log p_\theta(\mathbf{x}, \mathbf{z}) - \log q_\phi(\mathbf{z}|\mathbf{x})] \\ &= \mathbb{E}_{q_\phi(\mathbf{z}|\mathbf{x})} [\log p_\theta(\mathbf{x}|\mathbf{z})] - \text{KL}(q_\phi(\mathbf{z}|\mathbf{x}) || p_\theta(\mathbf{z}))\end{aligned}$$

Maximized to find the optimal θ and ϕ

q_ϕ should be such that data x is reconstructed well from \mathbf{z} (high log-lik)

q_ϕ should also be simple (close to the prior)

- The Reparametrization Trick is commonly used to optimize the ELBO
- Posterior is inferred only over \mathbf{z} , and usually only point estimate on θ

Here θ collectively denotes all the parameters of the prior and likelihood

Using the idea of “Amortized Inference” (next slide)

Here ϕ collectively denotes all the parameters that define the inference network

Stochastic Variational Inference

- Latent variable models need to infer the posterior $p(\mathbf{z}_n | \mathbf{x}_n)$ for each observation \mathbf{x}_n
- This can be slow if we have lots of observations because
 1. We need to iterate over each $p(\mathbf{z}_n | \mathbf{x}_n)$
 2. Learning the global parameters needs wait for step 1 to finish for all observations N
 3. $\mathcal{L}(\phi) = \mathbb{E}_{q_\phi(z)} [\log p(\mathbf{x}, \mathbf{z}) - \log_{q_\phi}(\mathbf{z})]$ summation over all data points \mathbf{x}_i is costly for large N
- One way to address this is via Stochastic VI
 - Use mini-batches of data and stochastic gradients to optimize the ELBO.

$$\widehat{\nabla_\phi \mathcal{L}} = \frac{N}{M} \sum_{i \in batch} \nabla_\phi \mathbb{E}_{q_\phi} [\log p(x_i, z_i) - \log_{q_\phi}(z_i)]$$

Amortized Inference

- Amortized inference is another appealing alternative (used in VAE and other LVMs too)
- No need to learn ϕ_n 's (one per data point) but just a single NN with params \mathbf{W}
 - This will be our “encoder network” for learning \mathbf{z}_n
 - Also very efficient to get $p(\mathbf{z}_* | \mathbf{x}_*)$ for a new data point \mathbf{x}_*

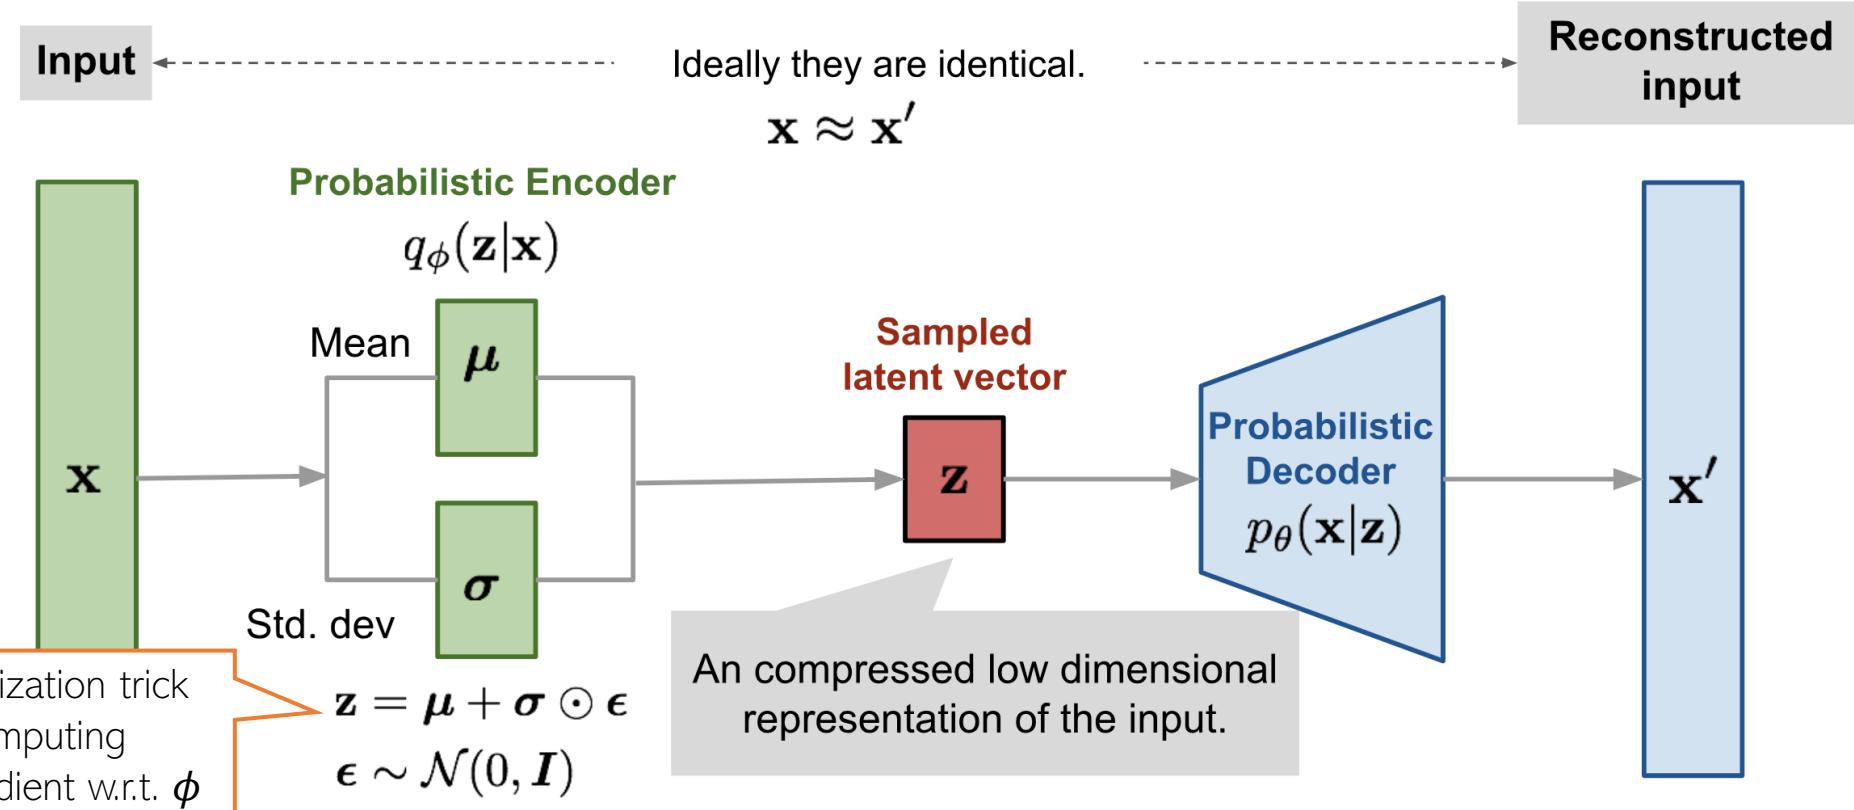
$$p(\mathbf{z}_n | \mathbf{x}_n) \approx q(\mathbf{z}_n | \phi_n) = q(\mathbf{z}_n | \text{NN}(\mathbf{x}_n; \mathbf{W}))$$

If q is Gaussian then the NN will output a mean and a variance

Variational Autoencoder: The Complete Pipeline

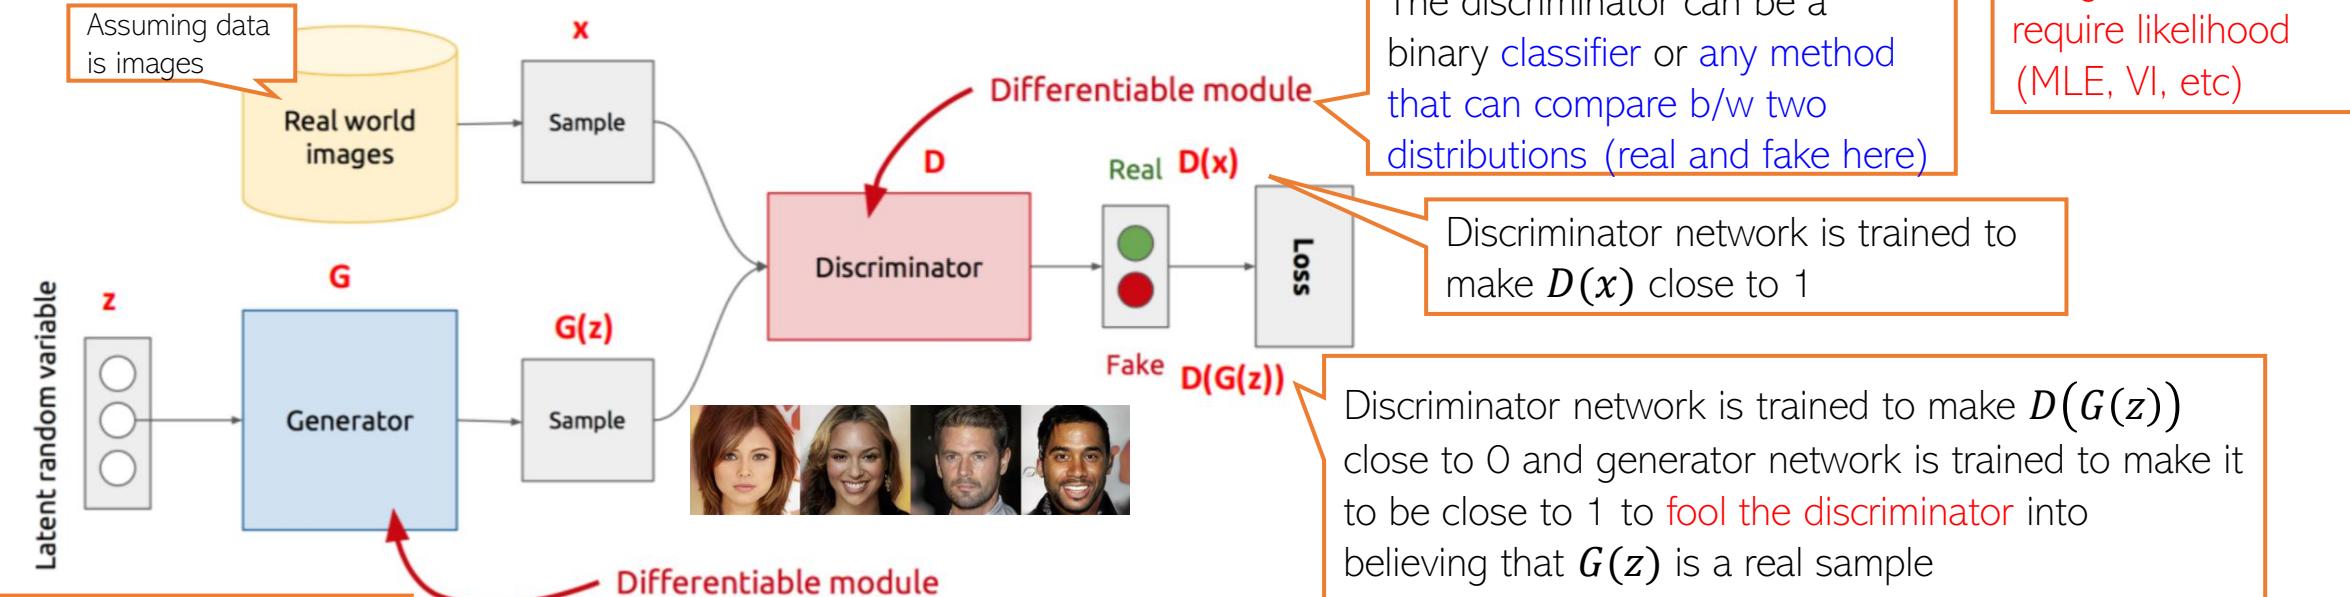
- Both probabilistic encoder and decoder learned jointly by maximizing the ELBO

$$\begin{aligned}\mathcal{L}(\theta, \phi | \mathbf{x}) &= \mathbb{E}_{q_\phi(\mathbf{z}|\mathbf{x})} [\log p_\theta(\mathbf{x}, \mathbf{z}) - \log q_\phi(\mathbf{z}|\mathbf{x})] \\ &= \mathbb{E}_{q_\phi(\mathbf{z}|\mathbf{x})} [\log p_\theta(\mathbf{x}|\mathbf{z})] - \text{KL}(q_\phi(\mathbf{z}|\mathbf{x}) \| p_\theta(\mathbf{z}))\end{aligned}$$



Generative Adversarial Network (GAN)

- GAN is an implicit generative latent variable model
- Can generate from it but can't compute $p(\mathbf{x})$ - the model doesn't define it explicitly
- GAN is trained using an **adversarial way** (Goodfellow et al, 2013)



$$\min_G \max_D V(D, G) = \mathbb{E}_{x \sim p_{data}(x)} [\log D(x)] + \mathbb{E}_{z \sim p_z(z)} [\log(1 - D(G(z)))]$$

Generative Adversarial Network (GAN)

- The GAN training criterion was

$$\min_G \max_D V(D, G) = \mathbb{E}_{x \sim p_{\text{data}}(x)} [\log D(x)] + \mathbb{E}_{z \sim p_z(z)} [\log(1 - D(G(z)))]$$

- With G fixed, the optimal D (exercise)

$$D_G^*(x) = \frac{p_{\text{data}}(x)}{p_{\text{data}}(x) + p_g(x)}$$

Distribution of real data

Distribution of synthetic data

- Given the optimal D , The optimal generator G is found by minimizing

$$\begin{aligned} V(D_G^*, G) &= \mathbb{E}_{x \sim p_{\text{data}}} \left[\log \frac{p_{\text{data}}(x)}{p_{\text{data}}(x) + p_g(x)} \right] + \mathbb{E}_{x \sim p_g} \left[\log \frac{p_g(x)}{p_{\text{data}}(x) + p_g(x)} \right] \\ &= \text{KL} \left[p_{\text{data}}(x) \middle\| \frac{p_{\text{data}}(x) + p_g(x)}{2} \right] + \text{KL} \left[p_g(x) \middle\| \frac{p_{\text{data}}(x) + p_g(x)}{2} \right] - \log 4 \end{aligned}$$

Jensen-Shannon divergence between p_{data} and p_g .
Minimized when $p_g = p_{\text{data}}$

Thus GAN can learn the true data distribution if the generator and discriminator have enough modeling power

GAN Optimization

$$\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{data}(\mathbf{x})}[\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_z(\mathbf{z})}[\log(1 - D(G(\mathbf{z})))]$$

- The GAN training procedure can be summarized as

- 1 Initialize θ_g, θ_d ; θ_g and θ_d denote the params of the deep neural nets defining the generator and discriminator, respectively
- 2 **for** each training iteration **do** In practice, for stable training, we run $K > 1$ steps of optimizing w.r.t. D and 1 step of optimizing w.r.t. G
- 3 **for** K steps **do**
- 4 Sample minibatch of M noise vectors $\mathbf{z}_m \sim q_z(\mathbf{z})$;
- 5 Sample minibatch of M examples $\mathbf{x}_m \sim p_D$;
- 6 Update the discriminator by performing stochastic gradient *ascent* using this gradient:

$$\nabla_{\theta_d} \frac{1}{M} \sum_{m=1}^M [\log D(\mathbf{x}_m) + \log(1 - D(G(\mathbf{z}_m)))].$$
 ;
- 7 Sample minibatch of M noise vectors $\mathbf{z}_m \sim q_z(\mathbf{z})$;
- 8 Update the generator by performing stochastic gradient *descent* using this gradient:

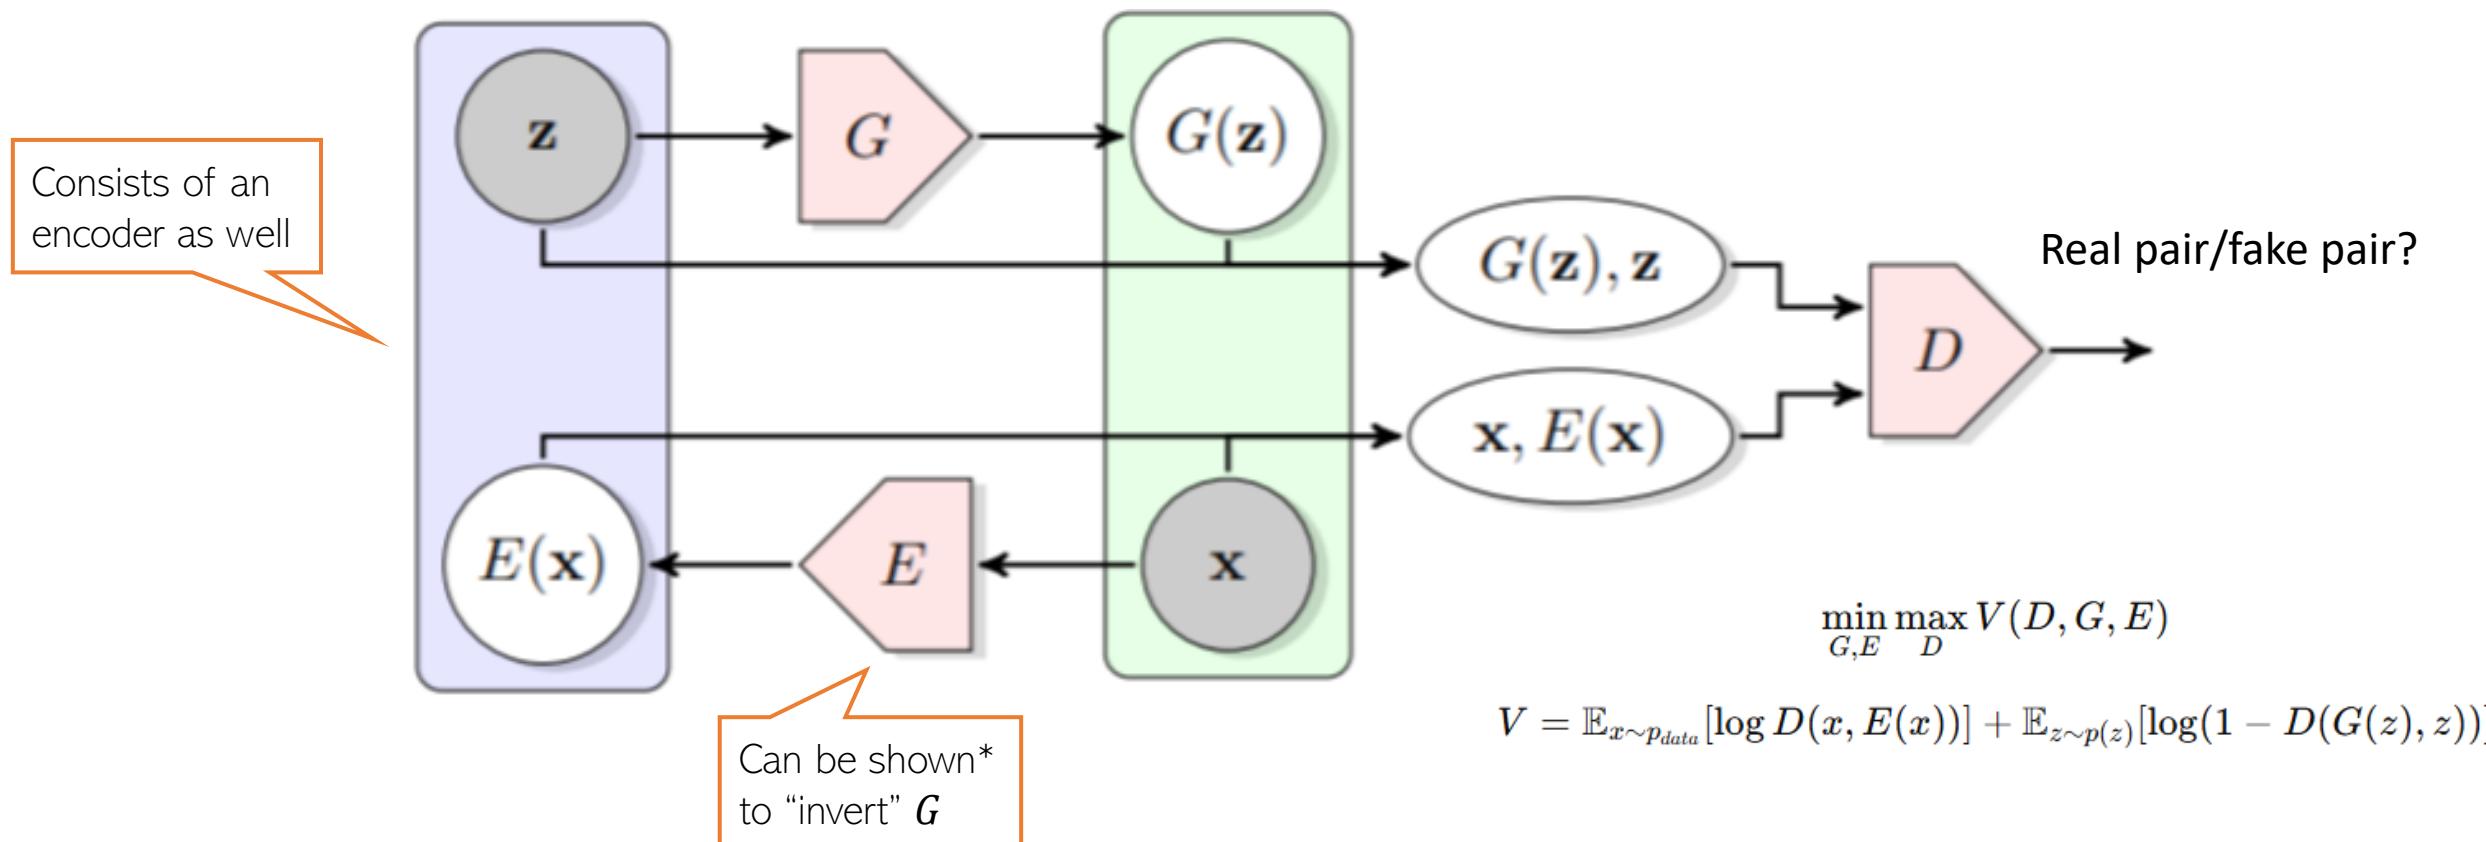
$$\nabla_{\theta_g} \frac{1}{M} \sum_{m=1}^M \log(1 - D(G(\mathbf{z}_m))).$$
 ;
- 9 Return θ_g, θ_d

In practice, in this step, instead of minimizing $\log(1 - D(G(z)))$, we **maximize** $\log(D(G(z)))$

Reason: Generator is bad initially so discriminator will always predict correctly initially and $\log(1 - D(G(z)))$ will saturate

GANs that also learn latent representations

- The standard GAN can only generate data. Can't learn the latent \mathbf{z} from \mathbf{x}
- Bidirectional GAN* (BiGAN) is a GAN variant that allows this



Evaluating GANs

- Two measures that are commonly used to evaluate GANs
 - Inception score (IS): Evaluates the distribution of generated data
 - Frechet inception distance (FID): Compares the distribution of real data and generated data
- Inception Score defined as $\exp(\mathbb{E}_{x \sim p_g} [\text{KL}(p(y|x) || p(y))])$ will be high if
 - Very few high-probability classes in each sample x : Low entropy for $p(y|x)$
 - We have diverse classes across samples: Marginal $p(y)$ is close to uniform (high entropy)
- FID uses extracted features (using a deep neural net) of real and generated data
 - Usually from the layers closer to the output layer
- These features are used to estimate two Gaussian distributions

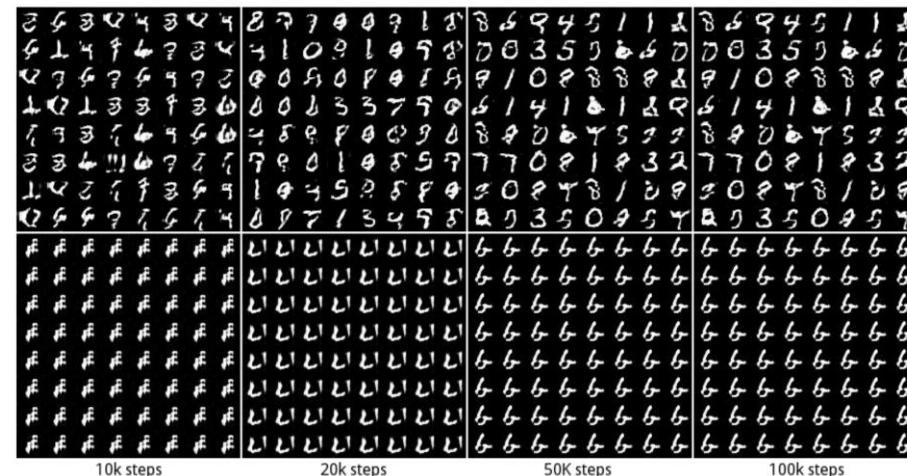
Using real data
 $\mathcal{N}(\mu_R, \Sigma_R)$
 $\mathcal{N}(\mu_G, \Sigma_G)$
Using generated data
- FID is then defined as $\text{FID} = |\mu_G - \mu_R|^2 + \text{trace}(\Sigma_G + \Sigma_R - (\Sigma_G \Sigma_R)^{1/2})$
- These measures can also be used for evaluating other deep gen models like VAE

High IS and low
FID is desirable

Both IS and FID measure how
realistic the generated data is

GAN: Some Issues/Comments

- GAN training can be hard and the basic GAN suffers from several issues
- Instability of training procedure
- Mode Collapse problem: Lack of diversity in generated samples
 - Generator may find some data that can easily fool the discriminator
 - It will stuck at that mode of the data distribution and keep generating data like that



GAN 1: No mode collapse (all 10 modes captured in generation)

GAN 2: Mode collapse (stuck on one of the modes)

- Some work on addressing these issues (e.g., [Wasserstein GAN](#), Least Squares GAN, etc)

References

- Chapter 21 and 26, Kevin Murphy, [Probabilistic Machine Learning: Advanced Topics](#)
- Chapter 19.2 and 17, Christopher N. Bishop, Deep Learning: Foundations and Concepts